

Working Notes for $D_{s1} \rightarrow D_s \gamma$ in Light-Cone QCD Sum Rules

Ds gamma project

July 5, 2026

Abstract

These notes are a living derivation for the light-cone QCD sum-rule calculation of $D_{s1}(2460) \rightarrow D_s \gamma$ and $D_{s1}(2536) \rightarrow D_s \gamma$. The purpose is pedagogical as well as practical: each step is written so that a master's student can follow the choices, conventions, and algebra before the final paper-level formulas are assembled.

1 Conventions

We use the mostly-minus metric

$$g_{\mu\nu} = \text{diag}(+1, -1, -1, -1), \quad (1)$$

and

$$\{\gamma_\mu, \gamma_\nu\} = 2g_{\mu\nu}, \quad \sigma_{\mu\nu} = \frac{i}{2}[\gamma_\mu, \gamma_\nu], \quad \gamma_5 = i\gamma^0\gamma^1\gamma^2\gamma^3. \quad (2)$$

The Levi-Civita convention is

$$\epsilon^{0123} = +1, \quad \epsilon_{0123} = -1, \quad (3)$$

so that

$$\text{Tr}[\gamma_5 \gamma_\mu \gamma_\nu \gamma_\alpha \gamma_\beta] = -4i \epsilon_{\mu\nu\alpha\beta}. \quad (4)$$

Our Fourier convention for a propagator is

$$S(x) = \int \frac{d^4k}{(2\pi)^4} e^{-ik \cdot x} \tilde{S}(k). \quad (5)$$

2 Currents and Momentum Assignment

The external momenta are

$$P = p + q, \quad (6)$$

where P is the initial axial-meson momentum, p is the final pseudoscalar D_s momentum, and q is the real-photon momentum.

The pseudoscalar interpolating current is

$$J_{D_s}(x) = \bar{s}(x) i\gamma_5 Q(x). \quad (7)$$

For the axial channel we use two independent currents,

$$J_{A\mu}^{(0)}(x) = \bar{s}(x) \gamma_\mu \gamma_5 Q(x), \quad (8)$$

$$J_{B\mu}^{(0)}(x) = \frac{i}{m_Q + m_s} \bar{s}(x) \sigma_{\mu\alpha} P^\alpha \gamma_5 Q(x). \quad (9)$$

The tensor current contracts with P^α , not p^α , because it interpolates the initial axial state. The factor $1/(m_Q + m_s)$ restores the same mass dimension as the axial-vector current.

3 Correlation Function

For a generic axial current $\Gamma_\mu \in \{\gamma_\mu \gamma_5, i\sigma_{\mu\alpha} P^\alpha \gamma_5 / (m_Q + m_s)\}$ we study

$$\Pi_\mu(p, q) = i \int d^4x e^{ip \cdot x} \langle \gamma(q, \varepsilon) | T \{ \bar{s}(x) \Gamma_\mu Q(x) \bar{Q}(0) i\gamma_5 s(0) \} | 0 \rangle. \quad (10)$$

The target Lorentz structure for the $1^+ \rightarrow 0^- \gamma$ transition is

$$(\eta \cdot \varepsilon^*)(p \cdot q) - (\eta \cdot q)(p \cdot \varepsilon^*). \quad (11)$$

On the correlator side we therefore project onto an equivalent transverse structure built from the open axial index and the photon polarization. With open indices this structure is represented by

$$T_{\mu\nu} = (p \cdot q)g_{\mu\nu} - q_\mu p_\nu, \quad S_{\mu\nu} \equiv -T_{\mu\nu} = p_\nu q_\mu - (p \cdot q)g_{\mu\nu}. \quad (12)$$

For a real photon, $q^2 = 0$, it obeys

$$q^\nu S_{\mu\nu} = 0, \quad (13)$$

so the projected invariant amplitude is gauge invariant under $\varepsilon_\nu \rightarrow \varepsilon_\nu + \lambda q_\nu$. In scalar form the same statement is

$$(\eta \cdot q)(p \cdot q) - (\eta \cdot q)(p \cdot q) = 0 \quad (14)$$

after replacing $\varepsilon \rightarrow q$.

4 Wick Contraction and OPE Bookkeeping

The OPE is separated into a hard-photon sector and a soft-photon sector:

$$\Pi_\mu^{\text{OPE}} = \Pi_\mu^{\text{hard}} + \Pi_\mu^{\text{soft}}. \quad (15)$$

This split is important because the two terms describe different physics and must not be double counted.

4.1 Hard photon emission

In the hard sector the photon is emitted perturbatively from either the heavy quark or the strange quark. After Wick contraction one obtains

$$\Pi_\mu^{\text{hard}} = iN_c \int d^4x e^{ip \cdot x} \text{Tr}_D \left[\Gamma_\mu S_Q^\gamma(x) i\gamma_5 S_s^0(-x) + \Gamma_\mu S_Q^0(x) i\gamma_5 S_s^\gamma(-x) \right], \quad (16)$$

where $N_c = 3$ and Tr_D denotes the Dirac trace. The free massive propagator is

$$S_q^0(x) = \int \frac{d^4k}{(2\pi)^4} e^{-ik \cdot x} \frac{i(\not{k} + m_q)}{k^2 - m_q^2}. \quad (17)$$

We keep explicit powers of m_s , including possible m_s^2 terms.

The first-order electromagnetic background-field correction for a generic quark flavor q is

$$S_q^\gamma(x) = -ie_q \int \frac{d^4k}{(2\pi)^4} e^{-ik \cdot x} \int_0^1 du \left[\frac{\not{k} + m_q}{2(m_q^2 - k^2)^2} F^{\alpha\beta}(ux) \sigma_{\alpha\beta} + \frac{ux_\alpha}{m_q^2 - k^2} F^{\alpha\beta}(ux) \gamma_\beta \right]. \quad (18)$$

In Eq. (16) this generic formula is used twice:

$$S_Q^\gamma(x) = S_q^\gamma(x)|_{q \rightarrow Q, m_q \rightarrow m_Q, e_q \rightarrow e_Q}, \quad (19)$$

$$S_s^\gamma(x) = S_q^\gamma(x)|_{q \rightarrow s, m_q \rightarrow m_s, e_q \rightarrow e_s}. \quad (20)$$

Thus the hard sector contains both heavy-line and strange-line electromagnetic background-field insertions. For a real photon we write

$$F^{\alpha\beta}(ux) = i(\varepsilon^\alpha q^\beta - \varepsilon^\beta q^\alpha) e^{iuq \cdot x}, \quad (21)$$

with the overall phase kept consistent with the Fourier convention.

Equation (18) is not a vacuum-condensate correction. It is the quark propagator in an external electromagnetic field and is included in the base analysis. What we do not include in the base hard sector are the ordinary SVZ-style condensate-expanded propagator terms such as $\langle \bar{s}s \rangle$, $\langle \bar{s}g_s \sigma G s \rangle$, or $\langle G^2 \rangle$ propagator corrections.

4.2 Equivalent momentum-space vertex form

For FeynCalc trace checks it is often simpler to use the equivalent explicit photon-vertex representation. With photon index ν ,

$$N_{Q,\mu\nu} = \text{Tr}_D[\Gamma_\mu(\not{k} + \not{q} + m_Q)\gamma_\nu(\not{k} + m_Q)i\gamma_5(\not{k} - \not{p} + m_s)], \quad (22)$$

$$N_{s,\mu\nu} = \text{Tr}_D[\Gamma_\mu(\not{k} + m_Q)i\gamma_5(\not{k} - \not{p} + m_s)\gamma_\nu(\not{k} - \not{p} + \not{q} + m_s)]. \quad (23)$$

These numerators are accompanied by the denominators

$$D_Q = [(k+q)^2 - m_Q^2][k^2 - m_Q^2][(k-p)^2 - m_s^2], \quad (24)$$

$$D_s = [k^2 - m_Q^2][(k-p)^2 - m_s^2][(k-p+q)^2 - m_s^2]. \quad (25)$$

The explicit-vertex form and the background-field form represent the same hard emission physics. We must not add both as independent contributions.

4.3 Soft photon emission: two-particle photon DAs

In the soft sector the photon is represented by its light-cone distribution amplitudes. Contracting only the heavy quark gives

$$\Pi_\mu^{\text{soft},2p} = i \int d^4x e^{ip \cdot x} [\Gamma_\mu S_Q^0(x) i\gamma_5]_{\alpha\beta} \langle \gamma(q, \varepsilon) | \bar{s}_\alpha(x) s_\beta(0) | 0 \rangle. \quad (26)$$

Fierz rearrangement converts the open spinor indices to a basis of Dirac bilinears,

$$\bar{s}_\alpha(x) s_\beta(0) = -\frac{1}{4} \sum_i (\Gamma_i)_{\beta\alpha} \bar{s}(x) \Gamma_i s(0), \quad (27)$$

so that

$$\Pi_\mu^{\text{soft},2p} = -\frac{i}{4} \int d^4x e^{ip \cdot x} \sum_i \text{Tr}_D[\Gamma_\mu S_Q^0(x) i\gamma_5 \Gamma_i] \langle \gamma(q, \varepsilon) | \bar{s}(x) \Gamma_i s(0) | 0 \rangle. \quad (28)$$

The photon matrix elements are then expressed in terms of two-particle photon DAs. Parameters such as $\chi \langle \bar{s}s \rangle$ enter here as part of the photon wavefunction normalization; they are not condensate corrections to the hard propagator.

4.4 Soft photon emission: three-particle photon DAs

To include three-particle photon DAs we must also allow one background gluon to come from the heavy-quark propagator,

$$S_Q^G(x) = -ig_s \int \frac{d^4k}{(2\pi)^4} e^{-ik \cdot x} \int_0^1 du \left[\frac{\not{k} + m_Q}{2(m_Q^2 - k^2)^2} G^{\alpha\beta}(ux) \sigma_{\alpha\beta} + \frac{ux_\alpha}{m_Q^2 - k^2} G^{\alpha\beta}(ux) \gamma_\beta \right]. \quad (29)$$

This term is not a gluon-condensate correction. It is an operator insertion that combines with photon matrix elements of the form

$$\langle \gamma(q, \varepsilon) | \bar{s}(x) \Gamma_i g_s G_{\alpha\beta}(ux) s(0) | 0 \rangle, \quad (30)$$

which define the three-particle photon DAs.

5 Staged Calculation Plan

We proceed in two stages.

Stage 1. Include the hard photon emission terms and the two-particle photon DAs:

$$\Pi_\mu^{\text{Stage 1}} = \Pi_\mu^{\text{hard}} + \Pi_\mu^{\text{soft},2p}. \quad (31)$$

This stage keeps all explicit m_s terms and omits ordinary condensate-expanded propagators.

Stage 2. Add the three-particle photon DA contributions generated by S_Q^G and the corresponding light-cone matrix elements:

$$\Pi_\mu^{\text{Stage 2}} = \Pi_\mu^{\text{Stage 1}} + \Pi_\mu^{\text{soft},3p}. \quad (32)$$

This staging keeps the derivation auditable: first we verify the core hard and two-particle DA structures, then we add the more involved background-gluon terms.

6 Current FeynCalc Artifacts

The script `scripts/step2_current_building_blocks.wl` checks the basic Dirac conventions and defines the current vertices. The script `scripts/step2_hard_photon_numerators.wl` computes numerator-level hard-photon traces in the explicit photon-vertex representation. These traces are stored in `outputs/step2_hard_photon_numerators.txt`.

The script `scripts/step2_e1_project_hard_numerators.wl` applies the E1 projector

$$\mathcal{P}_{E1}^{\mu\nu} = \frac{p^\nu q^\mu - (p \cdot q)g^{\mu\nu}}{2(p \cdot q)^2}, \quad (33)$$

which satisfies $\mathcal{P}_{E1}^{\mu\nu}[p_\nu q_\mu - (p \cdot q)g_{\mu\nu}] = 1$ for $q^2 = 0$. The projected hard numerators are stored in `outputs/step2_e1_projected_hard_numerators.txt`.

The companion script `scripts/ward_transversality_check.wl` verifies the Ward/transversality properties of the same E1 tensor with FeynCalc. Its output is stored in `outputs/ward_transversality_check`.

The script `scripts/step2_soft_2p_fierz_traces.wl` computes the trace factors multiplying the two-particle photon DA matrix elements after the Fierz rearrangement. The output `outputs/step2_soft_2p_fierz_traces.txt` shows which bilinears feed the axial current and which feed the tensor current.

7 Stage 1 Axial-Current Numerical Baseline

Before deriving the tensor-current sum rule, we implemented the known single-axial-current result for $D_{s1}(2460) \rightarrow D_s \gamma$ as a benchmark. This is the coupling called g_1 by Colangelo, De Fazio, and Ozpineci. In our notation this is the axial-current form factor G_A .

The Stage 1 version contains

$$G_A^{\text{Stage 1}} = G_A^{\text{pert}} + G_A^{\chi\phi\gamma} + G_A^{\mathbb{A},\bar{H}\gamma} + G_A^{f_{3\gamma}\Psi^v}, \quad (34)$$

where the three-particle DA integral usually denoted by $\mathcal{F}_1$ is omitted until Stage 2.

The numerical script is

`scripts/stage1_axial_g1_baseline.py`.

It writes the full Borel/threshold grid to

`outputs/stage1_axial_g1_baseline.csv`.

7.1 Magnetic-susceptibility sign convention

The Colangelo paper prints $\chi = -(3.15 \pm 0.30) \text{ GeV}^{-2}$ in its matrix-element convention. In this project we store the BBK magnitude as

$$\chi(1 \text{ GeV}) = +3.15 \pm 0.30 \text{ GeV}^{-2}. \quad (35)$$

Therefore the sign of the $\chi\phi_\gamma$ term in the implemented sum rule is flipped relative to the literal printed formula. This is a convention translation, not a physical change. A diagnostic run that inserts positive χ into the printed negative- χ formula gives a much larger width and is kept only as a warning check.

7.2 Preliminary Stage 1 result

Using the project inputs, $u_0 = 1/2$, $3 \leq M^2 \leq 6 \text{ GeV}^2$, and $s_0 = (2.50, 2.55, 2.60)^2 \text{ GeV}^2$, the axial-current Stage 1 grid gives

$$-0.399 \leq G_A^{\text{Stage } 1} \leq -0.314 \text{ GeV}^{-1}. \quad (36)$$

With

$$\Gamma = \frac{\alpha_{\text{em}}}{3} G^2 |\vec{q}_\gamma|^3, \quad (37)$$

this corresponds to

$$20.8 \text{ keV} \leq \Gamma_A^{\text{Stage } 1} [D_{s1}(2460) \rightarrow D_s \gamma] \leq 33.5 \text{ keV}. \quad (38)$$

This is not yet the physical mixed-state result. It is the axial-current baseline that will be combined with the tensor-current form factor after the J_B sum rule is derived.

8 First Tensor-Current Step: Leading-Twist Soft Contribution

The tensor-current calculation starts from the trace-level identity produced by `scripts/step3_soft_2p_e1_pro`. For the two-particle tensor photon DA matrix element, the E1 projections are

$$\mathcal{P}_{E1} [\text{Tr}(J_A S_Q i\gamma_5 \sigma_{\rho\lambda})(\varepsilon^\rho q^\lambda - \varepsilon^\lambda q^\rho)] = 8 \frac{k \cdot q}{p \cdot q}, \quad (39)$$

$$\mathcal{P}_{E1} [\text{Tr}(J_B S_Q i\gamma_5 \sigma_{\rho\lambda})(\varepsilon^\rho q^\lambda - \varepsilon^\lambda q^\rho)] = 8 \frac{m_Q}{m_Q + m_s}. \quad (40)$$

After the two-particle DA Fourier integral and double Borel transform, $k = p + \bar{u}q$, hence $k \cdot q = p \cdot q$ for $q^2 = 0$. Therefore, before the hadronic decay-constant normalization, the leading-twist tensor-current OPE piece is

$$\Pi_B^{\chi\phi_\gamma} = \frac{m_Q}{m_Q + m_s} \Pi_A^{\chi\phi_\gamma}. \quad (41)$$

The exploratory script `scripts/stage1_tensor_gb_leading_twist.py` uses this relation to estimate only the $\chi\phi_\gamma$ part of G_B . This is not the final tensor-current sum rule. It omits the tensor-current hard contribution, the vector DA contribution, and the Stage 2 three-particle DA terms.

Using the project inputs, the partial leading-twist estimate gives

$$G_B^{\chi\phi_\gamma} \simeq +0.205 \text{ to } +0.273 \text{ GeV}^{-1}, \quad \text{if } f_B = f_T, \quad (42)$$

$$G_B^{\chi\phi_\gamma} \simeq +0.114 \text{ to } +0.151 \text{ GeV}^{-1}, \quad \text{if } f_B = f_T \frac{m_{D_{s1}}}{m_c + m_s}. \quad (43)$$

The sizeable sign difference relative to G_A means that the physical mixed amplitudes can be sensitive to the tensor-current normalization. This is why the next step is to finish the full G_B sum rule before quoting final mixed widths.

9 Decay-Constant Normalization Before Mixing

The labels f_A , f_B , and f_T must not be mixed casually. The axial current couples to the axial basis state through

$$\langle 0 | J_A^\mu | D_{s1}^A(P, \eta) \rangle = f_A m_A \eta^\mu, \quad (44)$$

whereas our tensor-derived current is

$$J_B^\mu = \frac{i}{m_c + m_s} \bar{s} \sigma^{\mu\alpha} P_\alpha \gamma_5 c. \quad (45)$$

The corresponding decay constant is defined by

$$\langle 0 | J_B^\mu | D_{s1}^B(P, \eta) \rangle = f_B m_B \eta^\mu. \quad (46)$$

Therefore $f_B = f_T$ is not a physical statement by itself. It is only true if the tabulated f_T is already defined for the same current J_B^μ and with the same mass normalization. If instead the input table gives the more common tensor-current matrix element

$$\langle 0 | \bar{s} \sigma_{\mu\nu} \gamma_5 c | D_{s1}^B(P, \eta) \rangle = f_T (\eta_\mu P_\nu - \eta_\nu P_\mu), \quad (47)$$

then contraction with $P^\nu / (m_c + m_s)$ gives the effective normalization

$$f_B^{\text{eff}} \simeq f_T \frac{m_{D_{s1}}}{m_c + m_s}, \quad (48)$$

up to the sign fixed by the precise convention for $\sigma_{\mu\nu} \gamma_5$ and by the transverse condition $\eta \cdot P = 0$.

Only after G_A and G_B have been computed with their own basis-current decay constants do we form the physical amplitudes:

$$G_{2460} = \sin \theta G_A + \cos \theta G_B, \quad (49)$$

$$G_{2536} = \cos \theta G_A - \sin \theta G_B. \quad (50)$$

Thus the sine-cosine factors belong to the physical-state mixing of amplitudes, not to an arbitrary identification $f_B = f_T$. In the numerical scripts, $f_B = f_T$ is kept only as a diagnostic. The converted $f_B = f_T m_{D_{s1}} / (m_c + m_s)$ option is the safer central normalization unless the input-table definition of f_T is changed.

9.1 Physical-current two-point normalization check

There is one further normalization issue that should be separated from the three-point OPE calculation. The two basis currents J_A^μ and J_B^μ are not necessarily the same as the physical currents that interpolate only $D_{s1}(2460)$ or only $D_{s1}(2536)$. If we define

$$J_1^\mu = \sin \theta J_A^\mu + \cos \theta J_B^\mu, \quad (51)$$

$$J_2^\mu = \cos \theta J_A^\mu - \sin \theta J_B^\mu, \quad (52)$$

then the physical decay constants should be defined by

$$\langle 0 | J_1^\mu | D_{s1}(2460)(P, \eta) \rangle = f_1 M_{2460} \eta^\mu, \quad (53)$$

$$\langle 0 | J_2^\mu | D_{s1}(2536)(P, \eta) \rangle = f_2 M_{2536} \eta^\mu. \quad (54)$$

These constants are, in principle, obtained from two-point sum rules, not by identifying them directly with f_A and f_B . For the transverse invariant amplitudes of the two-point functions one has schematically

$$\Pi_{11} = \sin^2 \theta \Pi_{AA} + \cos^2 \theta \Pi_{BB} + 2 \sin \theta \cos \theta \Pi_{AB}, \quad (55)$$

$$\Pi_{22} = \cos^2 \theta \Pi_{AA} + \sin^2 \theta \Pi_{BB} - 2 \sin \theta \cos \theta \Pi_{AB}, \quad (56)$$

and after Borel transformation and continuum subtraction,

$$f_1^2 M_{2460}^2 \exp(-M_{2460}^2/M^2) = \hat{B} \Pi_{11}^{\text{OPE}}, \quad (57)$$

$$f_2^2 M_{2536}^2 \exp(-M_{2536}^2/M^2) = \hat{B} \Pi_{22}^{\text{OPE}}. \quad (58)$$

The mixed off-diagonal invariant

$$\Pi_{12} = \sin \theta \cos \theta (\Pi_{AA} - \Pi_{BB}) + (\cos^2 \theta - \sin^2 \theta) \Pi_{AB} \quad (59)$$

is a useful diagnostic: for an exactly diagonal physical-current basis it should be small in the chosen Borel window, or it can be used to determine the preferred mixing angle.

This affects only the final hadronic normalization, not the hard-photon or photon-DA OPE kernels already derived in the three-point calculation. If the three-point invariant amplitudes for the basis currents are written as

$$\Pi_A^{(3pt)} \propto f_A G_A, \quad (60)$$

$$\Pi_B^{(3pt)} \propto f_B G_B, \quad (61)$$

then the physical-current normalized couplings are

$$G_{2460}^{\text{phys}} = \frac{\sin \theta f_A G_A + \cos \theta f_B G_B}{f_1}, \quad (62)$$

$$G_{2536}^{\text{phys}} = \frac{\cos \theta f_A G_A - \sin \theta f_B G_B}{f_2}, \quad (63)$$

up to the obvious mass factors if the two basis-current residues are defined with different hadron-mass normalizations. The numerical tables above instead use the basis-current-normalized combination

$$G_{2460}^{\text{basis}} = \sin \theta G_A + \cos \theta G_B, \quad (64)$$

$$G_{2536}^{\text{basis}} = \cos \theta G_A - \sin \theta G_B. \quad (65)$$

Therefore the existing results should be kept, but labeled as the “basis-current normalization baseline”. They are not invalid: they contain the completed Stage-1 and Stage-2 OPE calculation, and they are the correct starting point for phenomenology. A final physical-current normalization update requires evaluating the two-point sum rules for f_1 and f_2 using Π_{AA} , Π_{BB} and Π_{AB} . If f_1 and f_2 are close to the effective residues assumed in the basis normalization, the present tables will change only mildly; if not, the couplings in Eqs. (62)–(63) should replace the basis-current values in the final paper tables.

10 Tensor-Current Hard Term: Triangle Reduction

The hard tensor-current numerator is more complicated than the axial one because the J_B current introduces an extra factor of P^α and therefore higher powers of the loop momentum after E1 projection. The script `scripts/step3_tensor_hard_triangle_reduction.wl` rewrites scalar products such as k^2 , $k \cdot p$, and $k \cdot q$ in terms of inverse propagator denominators. For the heavy-line photon emission we define

$$d_{H1} = (k + q)^2 - m_Q^2, \quad (66)$$

$$d_{H2} = k^2 - m_Q^2, \quad (67)$$

$$d_{H3} = (k - p)^2 - m_s^2, \quad (68)$$

and similarly for the strange-line photon emission. Setting $d_{Hi} = 0$ isolates the genuine three-denominator triangle core. For the tensor current this gives

$$N_{B,H}^{\text{core}} = \frac{-8im_Q^2 + 4im_Qm_s}{m_Q + m_s} - \frac{4im_Q^2p^2}{(m_Q + m_s)(p \cdot q)}, \quad (69)$$

$$N_{B,s}^{\text{core}} = \frac{-4im_Qm_s}{m_Q + m_s} - \frac{4im_s^2p^2}{(m_Q + m_s)(p \cdot q)}. \quad (70)$$

The same reduction was also performed for the axial current in `scripts/step3_axial_hard_triangle_red`. This is a useful calibration because the axial hard term has a known published spectral density.

10.1 Why denominator-cancellation terms are separated

Terms proportional to one or more inverse denominators cancel one or more propagators. After such a cancellation the loop integral contains at most two propagators. These reduced integrals depend on at most one of the two external virtualities, p^2 or $(p + q)^2$, and therefore do not generate the double-pole structure selected by the double Borel transform. They are recorded in the output files for auditability, but the double-pole sum rule is built from the three-denominator triangle core.

This observation narrows the remaining hard-tensor task: derive the double spectral density associated with the tensor-current triangle core and match it to the same E1 invariant amplitude used for the axial-current benchmark.

11 First Hard-Tensor Spectral Density Candidate

The next step is to translate the triangle core into a perturbative spectral density. We do this by calibrating against the known axial-current result. For one photon-emitting line, let m_i be the emitting-quark mass and m_j the spectator mass. The axial triangle core has the schematic form

$$N_i^{A,\text{core}} = -4i \left[a_i + \frac{b_i}{p \cdot q} \right]. \quad (71)$$

Comparing with the published axial spectral density fixes the replacement

$$-4i a_i \longrightarrow 2a_i \log \frac{s - m_j^2 + m_i^2 - \lambda^{1/2}(s, m_j^2, m_i^2)}{s - m_j^2 + m_i^2 + \lambda^{1/2}(s, m_j^2, m_i^2)}, \quad (72)$$

$$-4i \frac{b_i}{p \cdot q} \longrightarrow \frac{b_i}{m_i^2} \frac{m_j^2 - m_i^2 - s}{s^2} \lambda^{1/2}(s, m_j^2, m_i^2). \quad (73)$$

The overall factor is $-3/(8\pi^2)$, and the total axial density is assembled as the strange-emission contribution minus the heavy-emission contribution, with the corresponding charges. The script

`scripts/stage1_tensor_gb_hard_candidate.py`

first applies these rules back to the axial current and reproduces Colangelo's $\rho^P(s)$ with a maximum numerical difference 1.4×10^{-17} on the test grid.

For the tensor current the core contains $p^2/(p \cdot q)$:

$$N_{B,H}^{\text{core}} = -4i \left[\frac{2m_Q^2 - m_Qm_s}{m_Q + m_s} + \frac{m_Q^2p^2}{(m_Q + m_s)(p \cdot q)} \right], \quad (74)$$

$$N_{B,s}^{\text{core}} = -4i \left[\frac{m_Qm_s}{m_Q + m_s} + \frac{m_s^2p^2}{(m_Q + m_s)(p \cdot q)} \right]. \quad (75)$$

The diagonal double-dispersion representation explains how the p^2 numerator is treated. For the double-pole part one has

$$\frac{p^2}{(s-p^2)(s-P^2)} = \frac{s}{(s-p^2)(s-P^2)} - \frac{1}{s-P^2}, \quad P = p+q. \quad (76)$$

The second term has only a single pole and is removed by the double Borel projection in p^2 and P^2 . Therefore, inside the double-pole sum rule, p^2 in the hard numerator can be replaced by the spectral variable s . The same argument gives $P^2 \rightarrow s$ for any numerator factor of P^2 . The script

`scripts/step4_double_pole_p2_reduction.py`

records this identity as a short audit check.

Combining this hard candidate with the tensor-DA two-particle soft contribution gives, over the same M^2 and s_0 windows,

$$G_B^{\text{hard}} \simeq -0.831 \text{ to } -0.673 \text{ GeV}^{-1}, \quad G_B^{\text{Stage } 1} \simeq -0.510 \text{ to } -0.440 \text{ GeV}^{-1}, \quad (f_B = f_T), \quad (77)$$

$$G_B^{\text{hard}} \simeq -0.460 \text{ to } -0.373 \text{ GeV}^{-1}, \quad G_B^{\text{Stage } 1} \simeq -0.283 \text{ to } -0.244 \text{ GeV}^{-1}, \quad \left(f_B = f_T \frac{m_{D_{s1}}}{m_c + m_s} \right). \quad (78)$$

At $\theta = 35.3^\circ$, the corresponding candidate mixed widths are

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] \simeq 62.5 \text{ to } 87.9 \text{ keV}, \quad \Gamma[D_{s1}(2536) \rightarrow D_s \gamma] \simeq 0.00 \text{ to } 0.30 \text{ keV}, \quad (f_B = f_T), \quad (79)$$

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] \simeq 30.9 \text{ to } 44.7 \text{ keV}, \quad \Gamma[D_{s1}(2536) \rightarrow D_s \gamma] \simeq 4.0 \text{ to } 8.2 \text{ keV}, \quad \left(f_B = f_T \frac{m_{D_{s1}}}{m_c + m_s} \right). \quad (80)$$

These numbers should be read as a controlled Stage 1 candidate, not yet as the final result, because Stage 2 three-particle photon DAs remain. The $p^2 \rightarrow s$ step is justified for the double-pole part; any single-pole remnants belong outside the double-Borel sum rule.

12 Stage 2 Axial Three-Particle DA Check

Before deriving the tensor-current three-particle terms, we implemented the known axial-current correction as a Stage 2 calibration. The published axial sum rule contains the combination

$$\mathcal{F}_1 = \mathcal{S} + \tilde{\mathcal{S}} - \mathcal{T}_1 - \mathcal{T}_2 + \mathcal{T}_3 + \mathcal{T}_4 + 2v(-\mathcal{S} - \mathcal{T}_3 + \mathcal{T}_2), \quad (81)$$

integrated over the two standard regions

$$0 \leq v \leq 1 - u_0, \quad 0 \leq \alpha_g \leq \frac{u_0}{1-v}, \quad (82)$$

$$1 - u_0 \leq v \leq 1, \quad 0 \leq \alpha_g \leq \frac{1-u_0}{v}, \quad (83)$$

with

$$\alpha_q = u_0 - (1-v)\alpha_g, \quad \alpha_{\bar{q}} = 1 - u_0 - v\alpha_g. \quad (84)$$

The script

`scripts/stage2_axial_g1_three_particle.py`

evaluates this integral using the three-particle photon DAs in `photon_da.py`. For $u_0 = 1/2$ it finds

$$\int \mathcal{F}_1 = +7.033 \times 10^{-2}, \quad (85)$$

$$\Delta G_A^{\mathcal{F}_1} \simeq +0.0012 \text{ to } +0.0020 \text{ GeV}^{-1}. \quad (86)$$

Thus the axial Stage 2 result over the same Borel and threshold window is

$$-0.398 \leq G_A^{\text{Stage } 2} \leq -0.312 \quad \text{GeV}^{-1}, \quad (87)$$

corresponding to

$$20.5 \text{ keV} \leq \Gamma_A^{\text{Stage } 2}[D_{s1}(2460) \rightarrow D_s \gamma] \leq 33.3 \text{ keV}. \quad (88)$$

This small correction is a useful calibration, but it does not yet complete Stage 2 for the mixed physical amplitudes. The remaining Stage 2 task is to derive the tensor-current analog of the three-particle photon DA contribution from the background-gluon insertion and project it onto the same E1 invariant.

13 First Tensor-Current Stage 2 Trace Pass

For the tensor-current three-particle term we first separated the local $\sigma_{\alpha\beta}$ part of the heavy-quark background-gluon propagator from the explicitly $x_\alpha \gamma_\beta$ part. The script

`scripts/step5_soft_3p_trace_kernels.wl`

computes the E1-projected traces for the local piece against the $\mathcal{S}$, $\tilde{\mathcal{S}}$, $\mathcal{T}_1$, and $\mathcal{T}_2$ Lorentz structures. The nonzero local kernels are

$$K_A[\mathcal{S}] = 8 \frac{k \cdot q}{p \cdot q}, \quad K_B[\mathcal{S}] = 8 \frac{m_Q}{m_Q + m_s}, \quad (89)$$

$$K_A[\mathcal{T}_1] = -32i \frac{k \cdot q}{p \cdot q}, \quad K_B[\mathcal{T}_1] = -32i \frac{m_Q}{m_Q + m_s}, \quad (90)$$

$$K_A[\mathcal{T}_2] = +32i \frac{k \cdot q}{p \cdot q}, \quad K_B[\mathcal{T}_2] = +32i \frac{m_Q}{m_Q + m_s}. \quad (91)$$

The $\tilde{\mathcal{S}}$ kernel vanishes for this local part. After the three-particle light-cone Fourier integral and double Borel projection, $k \cdot q = p \cdot q$, so the local trace ratio is again

$$\frac{K_B}{K_A} = \frac{m_Q}{m_Q + m_s}. \quad (92)$$

This is the same ratio encountered in the two-particle tensor-DA family.

At this point it is useful to separate the published axial combination into the piece generated by the $\sigma_{\alpha\beta}$ part of the background-field propagator and the piece generated by the $x_\alpha G^{\alpha\beta} \gamma_\beta$ part:

$$\mathcal{F}_1^\sigma = \mathcal{S} + \tilde{\mathcal{S}} - \mathcal{T}_1 - \mathcal{T}_2 + \mathcal{T}_3 + \mathcal{T}_4, \quad (93)$$

$$\mathcal{F}_1^{x\gamma} = 2v(-\mathcal{S} - \mathcal{T}_3 + \mathcal{T}_2), \quad \mathcal{F}_1 = \mathcal{F}_1^\sigma + \mathcal{F}_1^{x\gamma}. \quad (94)$$

The numerical split, evaluated in

`scripts/step10_stage2_f1_split.py`,

is

$$\int \mathcal{F}_1^\sigma = -1.5625 \times 10^{-1}, \quad (95)$$

$$\int \mathcal{F}_1^{x\gamma} = +2.2658 \times 10^{-1}, \quad (96)$$

$$\int \mathcal{F}_1 = +7.0330 \times 10^{-2}. \quad (97)$$

Thus the small published axial Stage 2 correction is partly the result of a cancellation between two larger pieces. The simple trace ratio $m_Q/(m_Q + m_s)$ is established for the $\sigma_{\alpha\beta}$ part, so the controlled sigma-matched tensor-current correction is

$$\Delta G_B^{3p} \simeq -0.00366 \text{ to } -0.00210 \text{ GeV}^{-1}, \quad (f_B = f_T \text{ diagnostic}), \quad (98)$$

$$\Delta G_B^{3p} \simeq -0.00203 \text{ to } -0.00116 \text{ GeV}^{-1}, \quad \left(f_B = f_T \frac{m_{D_{s1}}}{m_c + m_s} \text{ central} \right). \quad (99)$$

The corresponding sigma-matched Stage 2 mixed-width estimate at $\theta = 35.3^\circ$ is

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] \simeq 62.9 \text{ to } 88.2 \text{ keV}, \quad \Gamma[D_{s1}(2536) \rightarrow D_s \gamma] \simeq 0.00 \text{ to } 0.26 \text{ keV}, \quad (f_B = f_T \text{ diagnostic}), \quad (100)$$

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] \simeq 30.9 \text{ to } 44.8 \text{ keV}, \quad \Gamma[D_{s1}(2536) \rightarrow D_s \gamma] \simeq 3.8 \text{ to } 8.0 \text{ keV}, \quad \left(f_B = f_T \frac{m_{D_{s1}}}{m_c + m_s} \text{ central} \right) \quad (101)$$

If, instead, one multiplies the full axial $\mathcal{F}_1$ integral by the same trace ratio, one obtains only a diagnostic estimate:

$$\Delta G_B^{3p} \simeq +0.00094 \text{ to } +0.00165 \text{ GeV}^{-1}, \quad (f_B = f_T \text{ diagnostic}), \quad (102)$$

$$\Delta G_B^{3p} \simeq +0.00052 \text{ to } +0.00091 \text{ GeV}^{-1}, \quad \left(f_B = f_T \frac{m_{D_{s1}}}{m_c + m_s} \text{ central} \right), \quad (103)$$

with central-normalization widths

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] \simeq 30.5 \text{ to } 44.5 \text{ keV}, \quad (104)$$

$$\Gamma[D_{s1}(2536) \rightarrow D_s \gamma] \simeq 4.0 \text{ to } 8.1 \text{ keV}. \quad (105)$$

The difference between the sigma-matched and full-ratio diagnostic numbers is small at the current precision, but conceptually important: the $x_\alpha \gamma_\beta$ part must be matched separately before the tensor Stage 2 correction is final.

14 Stage 2 Tensor: Explicit x -Dependent Kernels

The next audit step keeps the coordinate vector x explicit. For the $\sigma_{\alpha\beta}$ part of S_Q^G , the raw E1 projections of the $\mathcal{T}_3$ and $\mathcal{T}_4$ structures are

$$K_A[\mathcal{T}_3] = +16i \frac{k \cdot q}{p \cdot q}, \quad K_B[\mathcal{T}_3] = +16i \frac{m_Q}{m_Q + m_s}, \quad (106)$$

$$K_A[\mathcal{T}_4] = -16i \frac{k \cdot q}{p \cdot q}, \quad K_B[\mathcal{T}_4] = -16i \frac{m_Q}{m_Q + m_s}. \quad (107)$$

Thus these two structures also obey the same ratio $K_B/K_A = m_Q/(m_Q + m_s)$ after $k \cdot q = p \cdot q$. This check is recorded in

`scripts/step6_soft_3p_x_dependent_kernels.wl.`

The genuinely new complication comes from the second term in the background propagator,

$$\frac{u x_\alpha}{m_Q^2 - k^2} G^{\alpha\beta}(ux) \gamma_\beta. \quad (108)$$

The script

`scripts/step7_soft_3p_xgamma_kernels.wl`

computes the raw kernels with x kept explicit. Unlike the previous kernels, the tensor-to-axial ratio is not a constant at this stage. For example the $\mathcal{S}$ ratio after only $k \cdot q \rightarrow p \cdot q$ contains

$$\frac{(p \cdot q)(k \cdot x - p \cdot x) + (k \cdot p + p^2 + 2p \cdot q)(q \cdot x)}{2m_Q^2(q \cdot x)}, \quad (109)$$

and analogous expressions occur for $\mathcal{T}_1$, $\mathcal{T}_2$, and $\mathcal{T}_3$. Therefore the tensor Stage 2 three-particle result cannot yet be promoted from an estimate to a final term. The remaining calculation is to perform the Fourier/Borel replacement of the explicit $x_\alpha/(q \cdot x)$ factors, including the derivative action on the heavy-quark denominator, and then recombine the result with the photon DA structures.

As a diagnostic, we can impose only the light-cone momentum routing

$$k = p + aq, \quad a = \alpha_{\bar{q}} + v\alpha_g, \quad q^2 = 0, \quad (110)$$

without yet applying the derivative replacement. The script

`scripts/step8_xgamma_saddle_reduction.py`

reduces the raw ratios $K_B/[(m_Q/(m_Q + m_s))K_A]$ to

$$R_{\mathcal{S}} = \frac{p^2 + (a+1)p \cdot q}{m_Q^2}, \quad (111)$$

$$R_{\mathcal{T}_1} = \frac{p^2 + (3a+2)p \cdot q}{3m_Q^2}, \quad (112)$$

$$R_{\mathcal{T}_2} = \frac{p^2 + (2-a)p \cdot q}{m_Q^2}, \quad (113)$$

$$R_{\mathcal{T}_3} = \frac{p^2 + 2a p \cdot q}{m_Q^2}. \quad (114)$$

If one additionally uses the heavy-line on-shell relation $m_Q^2 = k^2 = p^2 + 2a p \cdot q$, only $R_{\mathcal{T}_3}$ becomes exactly one. The other ratios remain nontrivial. This confirms that the $x_\alpha \gamma_\beta$ term must be treated with the full Fourier/Borel derivative machinery. It is not controlled by the simple trace-ratio argument used for the $\sigma_{\alpha\beta}$ part.

15 Stage 2 Tensor: Derivative Reduction of $x_\alpha \gamma_\beta$

The next step applies the Fourier derivative rules to the raw $x_\alpha \gamma_\beta$ kernels. With phase

$$\exp[i(p + aq - k) \cdot x], \quad (115)$$

integration over x gives, up to a common phase convention,

$$p \cdot x \rightarrow -i p \cdot \partial_k, \quad (116)$$

$$q \cdot x \rightarrow -i q \cdot \partial_k, \quad (117)$$

$$k \cdot x F(k) \rightarrow -i \partial_{k_\mu} [k_\mu F(k)] = -i [4F(k) + k \cdot \partial_k F(k)]. \quad (118)$$

The script

scripts/step9_xgamma_derivative_reduction.wl

applies these rules to the scalar kernels multiplying $1/(m_Q^2 - k^2)$. After the light-cone routing $k = p + aq$, the reduced axial kernels are proportional to

$$A_S^{x\gamma} = \frac{-8m_Q}{[-m_Q^2 + p^2 + 2a p \cdot q]^2}, \quad (119)$$

$$A_{\mathcal{T}_1}^{x\gamma} = \frac{48im_Q}{[-m_Q^2 + p^2 + 2a p \cdot q]^2}, \quad (120)$$

$$A_{\mathcal{T}_2}^{x\gamma} = \frac{16im_Q}{[-m_Q^2 + p^2 + 2a p \cdot q]^2}, \quad (121)$$

$$A_{\mathcal{T}_3}^{x\gamma} = \frac{-16im_Q}{[-m_Q^2 + p^2 + 2a p \cdot q]^2}. \quad (122)$$

The corresponding tensor-current kernels are not obtained by a universal $m_Q/(m_Q + m_s)$ multiplication. The ratios after derivative reduction are

$$R_S^\partial = 1 + \frac{(1-a)p \cdot q}{m_Q^2}, \quad (123)$$

$$R_{\mathcal{T}_1}^\partial = \frac{2m_Q^2 + p^2 + (2-a)p \cdot q}{3m_Q^2}, \quad (124)$$

$$R_{\mathcal{T}_2}^\partial = -2 + \frac{3p^2 + (2+3a)p \cdot q}{m_Q^2}, \quad (125)$$

$$R_{\mathcal{T}_3}^\partial = \frac{2m_Q^2 - p^2 - 2a p \cdot q}{m_Q^2}. \quad (126)$$

All these reduced kernels carry the squared heavy-quark denominator. For the double-pole residue we drop numerator terms proportional to $D = -m_Q^2 + p^2 + 2a p \cdot q$, equivalently using $m_Q^2 = p^2 + 2a p \cdot q$ in the residue. In the three-particle integration domains,

$$a = \alpha_{\bar{q}} + v\alpha_g = 1 - u_0, \quad (127)$$

so for the symmetric Borel point $u_0 = 1/2$ this residue is evaluated at $a = 1/2$.

The script

scripts/step11_tensor_xgamma_ratio_integral.py

uses these residue ratios to match the $\mathcal{S}$, $\mathcal{T}_2$, and $\mathcal{T}_3$ pieces of the tensor-current $x_\alpha \gamma_\beta$ contribution to the known axial normalization

$$\mathcal{F}_1^{x\gamma} = 2v(-\mathcal{S} - \mathcal{T}_3 + \mathcal{T}_2). \quad (128)$$

With the heavy-line residue condition and physical $p \cdot q$ in the E1 invariant, the matched $\mathcal{S}$, $\mathcal{T}_2$, $\mathcal{T}_3$ integral is

$$\int \mathcal{F}_{1,B}^{x\gamma}[\mathcal{S}, \mathcal{T}_2, \mathcal{T}_3] = +2.5818 \times 10^{-1}. \quad (129)$$

The tensor-current derivative reduction also produces a $\mathcal{T}_4$ residue even though the axial E1 kernel has $A_{\mathcal{T}_4}^{x\gamma} = 0$. This is not an obstacle: its normalization can be fixed directly relative to the axial $\mathcal{T}_2$ unit. On the double-pole residue,

$$\frac{B_{\mathcal{T}_4}^{x\gamma}}{[m_Q/(m_Q + m_s)] A_{\mathcal{T}_2}^{x\gamma}} = -\frac{(1-a)p \cdot q}{m_Q^2}. \quad (130)$$

This gives the tensor-only contribution

$$\int \mathcal{F}_{1,B}^{x\gamma}[\mathcal{T}_4] = -4.2136 \times 10^{-2}, \quad (131)$$

and therefore

$$\int \mathcal{F}_{1,B}^{x\gamma} = +2.1605 \times 10^{-1}. \quad (132)$$

Together with the sigma-like part this gives the full matched tensor three-particle integral

$$\int \left(\mathcal{F}_{1,B}^{\sigma} + \mathcal{F}_{1,B}^{x\gamma} \right) = +5.9796 \times 10^{-2}. \quad (133)$$

For comparison, the axial total is $\int \mathcal{F}_1 = +7.0330 \times 10^{-2}$. A diagnostic calculation that instead inserts the physical hadron value $p^2 = m_{D_s}^2$ directly into the OPE residue ratios gives a much larger integral, +1.5072, so we do not use it as the working tensor estimate.

Numerically, for the central decay-constant normalization $f_B = f_T m_{D_{s1}} / (m_c + m_s)$, the full matched Stage 2 tensor-current result gives

$$\Delta G_B^{3p} = +0.00044 \text{ to } +0.00078 \text{ GeV}^{-1}, \quad (134)$$

$$G_B^{\text{Stage 2}} = -0.2821 \text{ to } -0.2436 \text{ GeV}^{-1}. \quad (135)$$

At $\theta = 35.3^\circ$ this corresponds to

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] = 30.6 \text{ to } 44.5 \text{ keV}, \quad (136)$$

$$\Gamma[D_{s1}(2536) \rightarrow D_s \gamma] = 4.0 \text{ to } 8.1 \text{ keV}. \quad (137)$$

The final matched result is close to the older full-ratio diagnostic, but the derivation is now explicit about which pieces come from the $\sigma_{\alpha\beta}$ part, which come from $x_\alpha G^{\alpha\beta} \gamma_\beta$, and how the tensor-only $\mathcal{T}_4$ term is normalized.

16 Final Stage 2 Uncertainty Scan

The script

`scripts/final_stage2_uncertainty_scan.py`

performs a first independent-input uncertainty scan around the completed matched Stage 2 baseline. Each random point samples the Borel parameter $M^2 \in [3, 6] \text{ GeV}^2$, the continuum threshold $\sqrt{s_0} \in [2.50, 2.60] \text{ GeV}$, and the numerical inputs listed in `inputs_table.csv`. The tensor-current decay constant is converted as

$$f_B = f_T \frac{m_{D_{s1}}}{m_c + m_s}. \quad (138)$$

The scan is not yet a correlated global error analysis; it is a transparent Monte Carlo error budget for the present sum rule. It is the preferred band to quote because it allows all inputs, M^2 , and s_0 to fluctuate together.

For fixed $\theta = 35.3^\circ$, using 500 random points, the scan gives

$$G[D_{s1}(2460) \rightarrow D_s \gamma] = -0.424_{-0.067}^{+0.065} \text{ GeV}^{-1}, \quad (139)$$

$$G[D_{s1}(2536) \rightarrow D_s \gamma] = -0.140_{-0.051}^{+0.047} \text{ GeV}^{-1}, \quad (140)$$

where the uncertainty indicates the 16–84 percentile interval. The corresponding width estimates are

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] = 37.9_{-10.7}^{+12.9} \text{ keV}, \quad (141)$$

$$\Gamma[D_{s1}(2536) \rightarrow D_s \gamma] = 6.1_{-3.4}^{+5.2} \text{ keV}. \quad (142)$$

If instead the mixing angle is sampled uniformly over $25^\circ \leq \theta \leq 45^\circ$, the same scan gives

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] = 36.9^{+12.6}_{-10.8} \text{ keV}, \quad (143)$$

$$\Gamma[D_{s1}(2536) \rightarrow D_s \gamma] = 6.0^{+7.6}_{-4.1} \text{ keV}. \quad (144)$$

The full scan table is written to

`outputs/final_stage2_uncertainty_scan.csv.`

For visualization the script

`scripts/uncertainty_publication_plots.py`

also fits Gaussian curves to the Monte Carlo width histograms using the sample mean and sample standard deviation. These fits are descriptive rather than the primary quoted uncertainties, because the width distribution is mildly right-tailed, especially for $D_{s1}(2536)$. For the fixed-angle ensemble the Gaussian summaries are

$$\Gamma_{2460} = 38.99 \pm 12.76 \text{ keV}, \quad (145)$$

$$\Gamma_{2536} = 7.39 \pm 5.67 \text{ keV}, \quad (146)$$

while for $25^\circ \leq \theta \leq 45^\circ$ they are

$$\Gamma_{2460} = 38.47 \pm 12.78 \text{ keV}, \quad (147)$$

$$\Gamma_{2536} = 7.79 \pm 6.76 \text{ keV}. \quad (148)$$

The corresponding publication histograms are shown in Figs. 1 and 2.




Figure 1: Monte Carlo width distributions for fixed $\theta = 35.3^\circ$. The black curves are Gaussian overlays using the sample mean and standard deviation; the dashed vertical lines mark the 16–84 percentile interval.

The diagnostic script

`scripts/uncertainty_budget_one_at_time.py`

then varies one input group at a time around the central point. This is not the band to quote as the final uncertainty; it explains which input groups drive the Monte Carlo spread.

The main lesson from Table 1 is that the $D_{s1}(2460)$ width is most sensitive to condensate/ χ inputs and decay constants, while the $D_{s1}(2536)$ width is particularly sensitive to the mixing angle because the physical amplitude contains a cancellation between the axial and tensor-current components. The same ranking is visualized in Fig. 3.




Figure 2: Monte Carlo width distributions when the mixing angle is sampled uniformly over $25^\circ \leq \theta \leq 45^\circ$. The black curves are Gaussian overlays using the sample mean and standard deviation; the dashed vertical lines mark the 16–84 percentile interval.

Varied group	Γ_{2460} (keV)	Γ_{2536} (keV)
$M^2, \sqrt{s_0}$	37.4 [33.0, 41.1]	5.95 [5.05, 7.00]
θ	37.3 [33.7, 39.6]	5.83 [2.40, 11.2]
quark and hadron masses	37.1 [35.1, 39.1]	6.01 [4.92, 7.27]
decay constants	37.7 [32.9, 42.4]	5.91 [3.48, 9.42]
condensates and χ	38.7 [29.2, 48.6]	6.40 [4.00, 9.11]
photon DA shape	37.2 [37.2, 37.2]	6.04 [6.04, 6.04]
all inputs, fixed θ	37.4 [26.8, 50.2]	5.50 [2.73, 11.2]
all inputs and θ	36.1 [25.7, 48.3]	6.28 [1.78, 14.9]

Table 1: One-at-a-time 16–84 percentile uncertainty budget with 500 random points per group. The photon-shape row is flat in the present matched reduction; it should be read as negligible sensitivity in this implementation, not as a general statement about photon DAs.




Figure 3: Ranked one-at-a-time uncertainty sources. The plotted quantity is half of the 16–84 percentile width for each group.

16.1 Alternative lattice normalization of the leading-twist photon DA

The dominant uncertainty in the legacy scan is the combination $\chi\langle\bar{s}s\rangle$ entering the leading-twist photon DA. The Rohrwild/BBK-style baseline treats χ and $\langle\bar{s}s\rangle$ as separate inputs. A recent lattice calculation by Bacchio, Becirevic, Gagliardi, and Sanfilippo gives the strange-quark normalization directly,

$$f_{\gamma,s}^\perp(2 \text{ GeV}) = -51.0(1.8) \text{ MeV}, \quad (149)$$

where

$$f_{\gamma,s}^\perp(\mu) = \chi_s(\mu)\langle\bar{s}s\rangle(\mu). \quad (150)$$

To compare with the working $\mu = 1 \text{ GeV}$ setup, the script

`scripts/lattice_photon_normalization_comparison.py`

uses a leading-order tensor-current running factor $f_{\gamma,s}^\perp(1 \text{ GeV})/f_{\gamma,s}^\perp(2 \text{ GeV}) = 1.08$, giving

$$f_{\gamma,s}^\perp(1 \text{ GeV}) = -55.1(1.9) \text{ MeV}. \quad (151)$$

Only the leading-twist normalization is replaced in this comparison; the higher-twist and three-particle terms that contain $\langle\bar{s}s\rangle$ remain explicit. This makes the comparison conservative and easy to audit.

Input scenario	$\Gamma_{2460} \text{ (keV)}$	$\Gamma_{2536} \text{ (keV)}$
Legacy $\chi\langle\bar{s}s\rangle$, fixed θ	36.7 [26.8, 50.2]	5.60 [2.56, 11.0]
Lattice $f_{\gamma,s}^\perp$, fixed θ	14.4 [11.4, 18.8]	0.68 [0.12, 1.88]
Legacy $\chi\langle\bar{s}s\rangle$, θ scan	36.3 [24.8, 50.2]	5.74 [1.86, 12.9]
Lattice $f_{\gamma,s}^\perp$, θ scan	14.6 [11.3, 18.6]	0.64 [0.06, 2.50]

Table 2: Comparison of the legacy photon-DA normalization with the lattice-normalized leading-twist photon DA. The intervals are 16–84 percentiles from 500 random points per ensemble.

The lattice-normalized scenario substantially reduces the uncertainty and also shifts the central widths downward. This should be presented as an important alternative input scenario rather than folded silently into the legacy result, because it changes the central leading-twist normalization. The corresponding histograms are shown in Fig. 4.

For direct comparison with the legacy Monte Carlo histograms in Figs. 1 and 2, the same script also produces Gaussian-overlay plots for the lattice-normalized scenario. For fixed $\theta = 35.3^\circ$, the descriptive Gaussian summaries are

$$\Gamma_{2460}^{\text{lat. } f_\gamma^\perp} = 14.91 \pm 3.69 \text{ keV}, \quad (152)$$

$$\Gamma_{2536}^{\text{lat. } f_\gamma^\perp} = 0.97 \pm 0.97 \text{ keV}. \quad (153)$$

For the mixing-angle scan they are

$$\Gamma_{2460}^{\text{lat. } f_\gamma^\perp} = 14.94 \pm 3.77 \text{ keV}, \quad (154)$$

$$\Gamma_{2536}^{\text{lat. } f_\gamma^\perp} = 1.21 \pm 1.45 \text{ keV}. \quad (155)$$

The $D_{s1}(2536)$ distribution is visibly non-Gaussian near zero because the amplitude is cancellation-driven, so the percentile intervals in Table 2 remain the safer quoted uncertainties.

Photon-DA normalization comparison




Figure 4: Fixed-angle Monte Carlo comparison between the legacy $\chi(\bar{s}s)$ treatment and the lattice-normalized $f_{\gamma,s}^\perp$ treatment.

Lattice-normalized Monte Carlo width distribution, $\theta = 35.3^\circ$




Figure 5: Lattice-normalized Monte Carlo width distributions for fixed $\theta = 35.3^\circ$. The black curves are Gaussian overlays using the sample mean and standard deviation; dashed vertical lines mark the 16–84 percentile interval.

Lattice-normalized Monte Carlo width distribution, $25^\circ \leq \theta \leq 45^\circ$




Figure 6: Lattice-normalized Monte Carlo width distributions when the mixing angle is sampled over $25^\circ \leq \theta \leq 45^\circ$.

16.2 Recommended D_{s1} results table

For the manuscript, the safest way to report the width predictions is not a Gaussian mean and standard deviation, but the median with asymmetric 16–84 percentile intervals. This is especially important for $D_{s1}(2536)$, where the width is positive definite and the amplitude can pass close to zero because of the axial–tensor cancellation. Table 4 now uses the Stage-3 physical-current residue normalization. In this first implementation the off-diagonal two-point overlap is calibrated from the working Stage-3 anchor $f_1 = 0.345 \pm 0.017$ GeV, and the same overlap then predicts f_2 . Numerically this gives $f_2 \simeq 0.38$ GeV and a negative overlap $\rho_{AB} \simeq -0.4$. The previous basis-current table remains useful as a pre-residue diagnostic, but it should no longer be used as the headline D_{s1} result.

As a separate verification check, we compared the mixed-current residue model with the pure axial-current local-QCDSR benchmark in Wang’s Table 1, $f_A(D_{s1}) = 0.245 \pm 0.017$ GeV. This benchmark is not treated as an authority for the physical mixed residues f_1 and f_2 , because it uses the pure current $\bar{Q}\gamma_\mu\gamma_5q$. Therefore the comparison must be made only in a limiting-current configuration where one of our physical currents becomes J_A . The model gives $f_1 = f_A$ at $\theta = 90^\circ$, where $J_1 = J_A$, and $f_2 = f_A$ at $\theta = 0^\circ$, where $J_2 = J_A$. Our working pure-axial input, $f_A = 0.225 \pm 0.025$ GeV, is close to Wang’s 0.245 ± 0.017 GeV and agrees within the combined uncertainty; the small central-value difference is naturally attributable to different quark masses, condensates, continuum thresholds, Borel windows, OPE truncations, and normalization conventions. At the physical angle $\theta = 35.3^\circ$, the Wang-like pure-axial target lies near the lower edge of the mixed-current residue range and requires a strongly negative overlap, $\rho_{AB} \simeq -1$ with our working $f_A = 0.225$ GeV, or $\rho_{AB} \simeq -0.96$ if the Wang value is also used as the basis f_A . By contrast, the headline Stage-3 anchor $f_1 = 0.345$ GeV corresponds to $\rho_{AB} \simeq -0.41$ and $f_2 \simeq 0.381$ GeV. The check is stored in `outputs/stage3_residue_benchmark_checks.csv`.

Case	Current/target	θ	ρ_{AB}	f_A	f_1	f_2
Wang benchmark	pure J_A two-point current	–	–	0.245 ± 0.017	–	–
our pure- J_1 axial limit	$J_1 = J_A$	90°	0	0.225	0.225	0.462
our pure- J_2 axial limit	$J_2 = J_A$	0°	0	0.225	0.462	0.225
physical angle, Wang-like target	target $f_1 = 0.245$	35.3°	–1.00	0.225	0.247	0.451
physical angle, Wang f_A	target $f_1 = f_A = 0.245$	35.3°	–0.957	0.245	0.245	0.462
current Stage-3 anchor	target $f_1 = 0.345$	35.3°	–0.408	0.225	0.345	0.381

Table 3: Comparison of Wang’s pure axial-current benchmark with limiting cases of the present mixed-current residue model. Wang’s number checks the pure J_A scale, so the appropriate comparison is the pure- J_1 axial limit at $\theta = 90^\circ$ or the pure- J_2 axial limit at $\theta = 0^\circ$. The physical mixed-current residues require assumptions about the AA/BB overlap and are therefore not directly determined by Wang’s pure axial-current benchmark. All decay constants are in GeV. A machine-readable version is `outputs/wang_pure_axial_residue_comparison_table.csv`.

There is also a scale-convention issue that should be stated explicitly. The photon-DA normalizations in the table are quoted at $\mu = 1$ GeV, whereas the quark masses used in the hard amplitudes are the usual $\overline{\text{MS}}$ inputs, e.g. $m_s(2 \text{ GeV})$ and $m_c(m_c)$. This is a standard phenomenological input convention, but it is not a fully matched single-scale setup. A strict scale-consistent variant would evolve m_s , $\langle \bar{s}s \rangle$, and either χ_s or $f_{\gamma,s}^\perp$ to a common factorization scale before rerunning the sum rule. We therefore flag the present rows as mixed-scale inputs rather than interpreting them as a separate scale-consistent prediction.

The legacy rows should be retained for comparison with the older photon-DA literature. The lattice-normalized rows are the more modern input scenario and give a much tighter prediction for $D_{s1}(2460)$, while $D_{s1}(2536)$ remains best described by an asymmetric interval because the central amplitude is cancellation-suppressed. The machine-readable version of Table 4 is

Scenario	Photon input	implied χ_s	Mixing input	Scale check	Γ_{2460} (keV)	Γ_{2536} (keV)
Legacy	$\chi_s(1 \text{ GeV}) = 3.15 \pm 0.30$ (Rohrwild/BBK)	3.15 [2.89, 3.47]	$\theta = 35.3^\circ$	mixed-scale input	$37.7^{+11.5}_{-10.5}$	$0.045^{+0.124}_{-0.040}$
Legacy	$\chi_s(1 \text{ GeV}) = 3.15 \pm 0.30$ (Rohrwild/BBK)	3.16 [2.86, 3.42]	$25^\circ \leq \theta \leq 45^\circ$	mixed-scale input	$37.8^{+10.3}_{-11.5}$	$0.350^{+0.752}_{-0.312}$
Lattice	$f_{\gamma,s}^\perp(1 \text{ GeV}) = -55.1 \pm 1.9 \text{ MeV}$ (Bacchio et al.)	4.99 [4.24, 6.05]	$\theta = 35.3^\circ$	mixed-scale input	$16.1^{+4.4}_{-3.7}$	$0.504^{+0.600}_{-0.377}$
Lattice	$f_{\gamma,s}^\perp(1 \text{ GeV}) = -55.1 \pm 1.9 \text{ MeV}$ (Bacchio et al.)	5.02 [4.25, 6.09]	$25^\circ \leq \theta \leq 45^\circ$	mixed-scale input	$15.8^{+5.0}_{-3.7}$	$0.498^{+1.087}_{-0.457}$

Table 4: Recommended reporting table for the $D_{s1} \rightarrow D_s \gamma$ widths. Entries are medians with 16–84 percentile intervals. The lattice scenario replaces only the leading-twist normalization $\chi_s \langle \bar{s}s \rangle$ by $f_{\gamma,s}^\perp$; higher-twist condensate terms are left explicit. The implied χ_s column is derived from $f_{\gamma,s}^\perp / \langle \bar{s}s \rangle$ in the lattice rows, so it inherits the strange-condensate spread. The scale-check column records that the present numerical rows use conventional mixed-scale phenomenological inputs. The widths include the Stage-3 physical-current residue correction calibrated with the working Stage-3 anchor $f_1 = 0.345 \pm 0.017 \text{ GeV}$; a fully matched-scale prediction should be generated as a separate rerun.

outputs/ds1_recommended_results_table.csv. The full Stage-3 residue-normalization scan is stored in outputs/stage3_ds1_physical_residue_scan.csv, with the compact summary in outputs/stage3_ds1_physical_residue_summary.csv.

17 Central Borel Stability

For the central input set and $\theta = 35.3^\circ$, the script

scripts/stage2_stability_plots.py

produces the Borel-stability curves shown in Fig. 7. The curves are smooth over $3 \leq M^2 \leq 6 \text{ GeV}^2$. At fixed $\sqrt{s_0}$, the dependence on M^2 is mild; changing $\sqrt{s_0}$ from 2.50 to 2.60 GeV shifts the normalization more visibly and should therefore remain part of the quoted systematic uncertainty.

Across the full displayed window the central ranges are

$$\Gamma[D_{s1}(2460) \rightarrow D_s \gamma] = 30.6 \text{ to } 44.5 \text{ keV}, \quad (156)$$

$$\Gamma[D_{s1}(2536) \rightarrow D_s \gamma] = 3.95 \text{ to } 8.10 \text{ keV}. \quad (157)$$

The corresponding amplitude spans are smaller:

$$G_{2460} = -0.460 \text{ to } -0.381 \text{ GeV}^{-1}, \quad (158)$$

$$G_{2536} = -0.162 \text{ to } -0.113 \text{ GeV}^{-1}. \quad (159)$$

This is the expected pattern: the widths amplify the residual sum-rule variation because $\Gamma \propto G^2 |\vec{q}|^3$.

18 Publication Figures and Phenomenology

The remaining publication-oriented figures are generated by

scripts/publication_plots.py.

Figure 8 shows the continuum-threshold dependence at fixed representative values of M^2 . It complements Fig. 7: the curves are monotonic in $\sqrt{s_0}$, so the continuum threshold should be quoted as an important systematic.

Figure 9 shows the mixing-angle dependence. The $D_{s1}(2460)$ width remains at the tens-of-keV level over a broad range of θ , while the $D_{s1}(2536)$ width is strongly suppressed near the




Figure 7: Central matched Stage 2 stability for $D_{s1}(2460) \rightarrow D_s \gamma$ and $D_{s1}(2536) \rightarrow D_s \gamma$ at $\theta = 35.3^\circ$. The three curves correspond to $\sqrt{s_0} = 2.50, 2.55, 2.60$ GeV.




Figure 8: Continuum-threshold dependence of the matched Stage 2 widths for $\theta = 35.3^\circ$ at three representative Borel points.




Figure 9: Mixing-angle dependence at $M^2 = 4.5 \text{ GeV}^2$ and $\sqrt{s_0} = 2.55 \text{ GeV}$. The dashed line marks $\theta = 35.3^\circ$.

large-angle region because its amplitude involves a cancellation between the axial and tensor-current components.

The cancellation pattern is displayed more explicitly in Fig. 10. At the central point $M^2 = 4.5 \text{ GeV}^2$, $\sqrt{s_0} = 2.55 \text{ GeV}$, the $D_{s1}(2536)$ amplitude receives opposite-sign contributions from the axial Stage 1 term and the hard tensor-current term. The three-particle DA corrections are numerically small in the final physical amplitudes, but they are included in the quoted Stage 2 result.




Figure 10: Contribution breakdown for the physical couplings at the central Borel and threshold point. The labels “A” and “B” refer to the axial-current and tensor-current components before physical mixing.

Finally, Fig. 11 compares the present result for $D_{s1}(2460) \rightarrow D_s \gamma$ with representative estimates collected by Colangelo, De Fazio and Ozpineci. The present result is higher than the original axial-current LCSR interval, but it remains the same order of magnitude. It should not be advertised as explaining a non-observation by making the $D_{s1}(2460) \rightarrow D_s \gamma$ width tiny. Instead, the robust phenomenological statement from this calculation is that the $D_{s1}(2536) \rightarrow D_s \gamma$

channel is much more cancellation-sensitive and can be substantially smaller.




Figure 11: Comparison of the present $D_{s1}(2460) \rightarrow D_s \gamma$ width with representative literature estimates. The Colangelo LCSR, VMD and quark-model values are from the comparison table in Colangelo, De Fazio and Ozpineci (hep-ph/0505195).

19 Validation Checks Before the B_{s1} Extension

The script

`scripts/validation_checks.py`

collects several checks that should be repeated whenever the numerical implementation is modified. The first check is an axial-current benchmark against Colangelo, De Fazio and Ozpineci, using paper-like input choices and the BBK value $\kappa = 0.2$ for the relevant photon DA parameter. Over $3 \leq M^2 \leq 6 \text{ GeV}^2$ and $2.50 \leq \sqrt{s_0} \leq 2.60 \text{ GeV}$, the present controlled benchmark gives

$$g_1 = -0.340 \text{ to } -0.237 \text{ GeV}^{-1}, \quad \Gamma = 11.8 \text{ to } 24.3 \text{ keV}. \quad (160)$$

The published comparison band is $g_1 = -0.37 \text{ to } -0.29 \text{ GeV}^{-1}$ and $\Gamma = 19 \text{ to } 29 \text{ keV}$. This is close enough in sign and scale to validate the implementation, but it also shows that an exact literature reproduction requires matching every convention and input choice. We should present this as a benchmark check, not as an independent final phenomenological result.

The same validation script checks the photon DA normalization and the three-particle integration closure:

$$\int_0^1 du \varphi_\gamma(u) = 1.0000000000, \quad (161)$$

$$\max_u |\psi^a(u) - \psi^a(1-u)| = 4.44 \times 10^{-15}, \quad (162)$$

$$\max_u |\psi^v(u) + \psi^v(1-u)| = 1.78 \times 10^{-15}, \quad (163)$$

$$F_1 - F_1^{(1)} - F_1^{(2)} = 0. \quad (164)$$

These are numerical quadrature checks of the DA implementation and the split integration domains.

The independent FeynCalc script `scripts/ward_transversality_check.wl` checks the selected E1 tensor $S_{\mu\nu} = p_\nu q_\mu - (p \cdot q)g_{\mu\nu}$. With $q^2 = 0$, $\varepsilon \cdot q = 0$, $P = p + q$, and $\eta \cdot P = 0$, it gives

$$S_{\mu\nu}q^\nu = 0, \quad (165)$$

$$[(\eta \cdot \varepsilon)(p \cdot q) - (\eta \cdot q)(p \cdot \varepsilon)]_{\varepsilon \rightarrow q} = 0, \quad (166)$$

$$P^\mu S_{\mu\nu} \varepsilon^\nu = 0, \quad (167)$$

$$\frac{S_{\mu\nu} S^{\mu\nu}}{2(p \cdot q)^2} = 1. \quad (168)$$

The first two equations are the photon Ward check, the third confirms compatibility with the physical initial axial-meson polarization, and the last line is the projector normalization.

At the central point $M^2 = 4.5 \text{ GeV}^2$ and $\sqrt{s_0} = 2.55 \text{ GeV}$, the hard tensor-current charge decomposition is

$$G_B^{\text{hard}} = -4.092 \times 10^{-1} \text{ GeV}^{-1}, \quad (169)$$

$$G_B^{\text{hard}}(e_c \text{ line}) = -2.499 \times 10^{-1} \text{ GeV}^{-1}, \quad (170)$$

$$G_B^{\text{hard}}(e_s \text{ line}) = -1.594 \times 10^{-1} \text{ GeV}^{-1}. \quad (171)$$

This confirms that hard photon emission from both quark lines is numerically important and must remain in the baseline.

Diagnostic case	G_A (GeV ⁻¹)	G_B	Γ_{2460} (keV)	Γ_{2536} (keV)
central	-0.357	-0.263	37.18	6.04
$m_s \rightarrow 0$ numerical proxy	-0.175	-0.210	15.59	0.15
$m_s = m_c$ toy	+0.535	+0.273	59.46	24.22
$e_c = 0$ toy	+0.102	-0.013	0.50	2.57
$e_s = 0$ toy	-0.459	-0.250	46.30	16.48
light- u OPE toy	-1.040	-0.438	192.96	110.02
light- d OPE toy	-0.109	-0.176	8.94	0.05

Table 5: Limiting-case diagnostics from `outputs/validation_limiting_cases.csv`. The $m_s \rightarrow 0$ and light-quark rows use $m_q = 10^{-4} \text{ GeV}$ as a numerical regulator of the present hard-line spectral-density representation. The light- u and light- d rows are OPE-level charge and condensate substitutions only, not physical nonstrange $D_1 \rightarrow D\gamma$ predictions.

20 First $B_{s1} \rightarrow B_s \gamma$ Extension

The same matched Stage-2 machinery can be reused for the bottom-strange system after the replacement $m_c \rightarrow m_b$ and $e_c \rightarrow e_b = -1/3$. The script

`scripts/bs1_stage2_baseline.py`

performs this first pass. Internally the reusable functions still use the key `mc`; in this section it should be read as m_b .

The first-pass numerical inputs are deliberately conservative working inputs: $m_b(m_b) = 4.18 \text{ GeV}$, $m_s(2 \text{ GeV}) = 0.093 \text{ GeV}$, $m_{B_s} = 5.36692 \text{ GeV}$, and $f_{B_s} = 0.2303 \text{ GeV}$ (FLAG Review 2024). For the axial and tensor B_{s1} decay constants we use the temporary heavy-quark scaling estimate $f_H \sqrt{m_H} = \text{const.}$ from the D_{s1} central values. Thus these rows are not yet final paper predictions; they are a controlled baseline for the bottom-sector calculation.

Two mass targets are shown. The observed narrow state is treated as the high-combination state,

$$G_{\text{high}} = \cos \theta G_A - \sin \theta G_B, \quad m_{B_{s1}} = 5.82870 \text{ GeV}. \quad (172)$$

Because the lower 1^+ bottom-strange partner is not firmly established, the second row uses $m_{B_{s1}} = 5.720 \text{ GeV}$ only as a lower-partner diagnostic, with

$$G_{\text{low}} = \sin \theta G_A + \cos \theta G_B. \quad (173)$$

The scan uses $8 \leq M^2 \leq 14 \text{ GeV}^2$, $\theta = 35.3^\circ$, and $\sqrt{s_0} = 6.05, 6.15, 6.25 \text{ GeV}$ for the observed state (5.95, 6.05, 6.15 GeV for the lower diagnostic state).

Target	Photon input	$G \text{ (GeV}^{-1}\text{)}$	$\Gamma \text{ (keV)}$
$B_{s1}(5830)$ high combo	legacy χ_s (Rohrwild/BBK)	0.0358 [0.0218, 0.0688]	0.27 [0.10, 1.01]
$B_{s1}(5830)$ high combo	lattice $f_{\gamma,s}^\perp$ (Bacchio et al.)	0.1496 [0.119, 0.216]	4.75 [3.02, 9.88]
lower 1^+ diagnostic	legacy χ_s (Rohrwild/BBK)	0.589 [0.502, 0.761]	33.8 [24.6, 56.5]
lower 1^+ diagnostic	lattice $f_{\gamma,s}^\perp$ (Bacchio et al.)	0.847 [0.725, 1.09]	69.9 [51.2, 116]

Table 6: First matched Stage-2 $B_{s1} \rightarrow B_s \gamma$ baseline. Brackets show the range over the Borel and continuum-threshold grid, not a full uncertainty interval. The B_{s1} decay constants are working heavy-quark-scaled inputs, so these values should be refined after dedicated B_{s1} hadronic inputs are fixed.

The qualitative behavior mirrors the D_{s1} calculation: the observed high-combination state is cancellation-sensitive, while the lower-combination diagnostic is much larger. The difference between the legacy and lattice photon-normalization rows is again substantial, so the bottom-sector analysis should keep both scenarios until the input convention is finalized. The machine-readable grid is `outputs/bs1_stage2_baseline.csv`.

The companion script

`scripts/bs1_uncertainty_and_plots.py`

performs a 500-point Monte Carlo scan for each state, photon-input scenario, and mixing-angle treatment. The observed $B_{s1}(5830)$ row samples $8 \leq M^2 \leq 14 \text{ GeV}^2$, $6.05 \leq \sqrt{s_0} \leq 6.25 \text{ GeV}$, the same quark and photon-DA inputs used in the D_{s1} scan, and assigns a conservative 20% width to the temporary heavy-quark-scaled B_{s1} axial and tensor decay constants. The lower-state diagnostic uses the lattice-inspired model mass $m_{B_{s1}} = 5.750 \pm 0.026 \text{ GeV}$ from Lang, Mohler, Prelovsek and Woloshyn.

After the physical-current normalization discussion above, the temporary HQ-scaled B_{s1} constants are not the preferred bottom-sector input. The first bottom-sector update used the script

`scripts/stage3_bs1_pz_decay_constants.py`

to rerun the same B_{s1} Monte Carlo scan with the Pullin–Zwicky heavy-light sum-rule inputs $f_{B_{s1}} = 0.305 \pm 0.027 \text{ GeV}$ and $f_{B_{s1}}^T = 0.285 \pm 0.027 \text{ GeV}$. These numbers should be interpreted as basis-current inputs, not directly as the final physical-current decay constants. The follow-up script

`scripts/stage3_bs1_physical_decay_constants.py`

uses the same basis-current inputs and closes the two-point matrix by requiring the physical currents to be diagonal, $\Pi_{12} = 0$. In this closure model

$$f_1^2 = s_\theta^2 f_A^2 + c_\theta^2 f_B^2 + 2s_\theta c_\theta \rho_{AB} f_A f_B, \quad f_2^2 = c_\theta^2 f_A^2 + s_\theta^2 f_B^2 - 2s_\theta c_\theta \rho_{AB} f_A f_B,$$

where

$$\rho_{AB} = -\frac{s_\theta c_\theta (f_A^2 - f_B^2)}{(c_\theta^2 - s_\theta^2) f_A f_B}.$$

Here $f_A = f_{B_{s1},A}^{\text{PZ}}$ and $f_B = f_{B_{s1}}^T m_{B_{s1}} / (m_b + m_s)$ is the effective tensor-current normalization in the same dimension as f_A . Samples with $|\rho_{AB}| > 1$ or negative f_i^2 are rejected. This is still a model closure rather than an independent local OPE derivation of Π_{AA} , Π_{BB} and Π_{AB} , but it does calculate the bottom physical-current f_1 and f_2 from the same mixed-current logic used for the charm normalization. The older PZ-only rows are retained as a basis-normalization comparison.

The current bottom-sector physical-current results are:

Target	Photon input	Mixing input	Physical constants	ρ_{AB}	g (GeV ⁻¹)	Γ (keV)
$B_{s1}(5830)$ high combo	legacy χ_s (Rohrwild/BBK)	$\theta = 35.3^\circ$	$f_1 = 0.436[0.380, 0.486]$, $f_2 = 0.218[0.138, 0.292]$	$+0.587[+0.242, +0.835]$	$-0.0793[-0.130, -0.0537]$	$1.33[0.612, 3.60]$
$B_{s1}(5830)$ high combo	lattice $f_{\gamma,s}^\perp$ (Bacchio et al.)	$\theta = 35.3^\circ$	$f_1 = 0.443[0.384, 0.492]$, $f_2 = 0.203[0.121, 0.291]$	$+0.631[+0.276, +0.866]$	$-0.0401[-0.0695, -0.0274]$	$0.342[0.160, 1.026]$
lower 1^+ diagnostic	legacy χ_s (Rohrwild/BBK)	$\theta = 35.3^\circ$	$f_1 = 0.428[0.367, 0.479]$, $f_2 = 0.226[0.126, 0.303]$	$+0.529[+0.215, +0.854]$	$+0.284[+0.227, +0.360]$	$9.50[5.97, 17.0]$
lower 1^+ diagnostic	lattice $f_{\gamma,s}^\perp$ (Bacchio et al.)	$\theta = 35.3^\circ$	$f_1 = 0.440[0.381, 0.490]$, $f_2 = 0.216[0.132, 0.285]$	$+0.584[+0.286, +0.849]$	$+0.391[+0.331, +0.457]$	$18.4[12.5, 28.1]$

Table 7: Bottom-strange physical-current normalization from the mixed-current diagonalization closure. Entries are medians with 16–84 percentile intervals. The Pullin–Zwicky quantities are used as the basis inputs f_A and f_T , while the table reports the derived physical-current f_1 and f_2 .

For reference, the basis/PZ comparison before the physical-current diagonalization was:

Target	Photon input	Mixing input	Decay constants	G (GeV ⁻¹)	Γ (keV)
$B_{s1}(5830)$ high combo	legacy χ_s (Rohrwild/BBK)	$\theta = 35.3^\circ$	PZ	$-0.0140[-0.0387, +0.0155]$	$0.109[0.008, 0.413]$
$B_{s1}(5830)$ high combo	legacy χ_s (Rohrwild/BBK)	$25^\circ \leq \theta \leq 45^\circ$	PZ	$-0.0126[-0.0527, +0.0364]$	$0.249[0.023, 0.819]$
$B_{s1}(5830)$ high combo	lattice $f_{\gamma,s}^\perp$ (Bacchio et al.)	$\theta = 35.3^\circ$	PZ	$+0.0356[+0.0029, +0.0694]$	$0.272[0.028, 1.024]$
$B_{s1}(5830)$ high combo	lattice $f_{\gamma,s}^\perp$ (Bacchio et al.)	$25^\circ \leq \theta \leq 45^\circ$	PZ	$+0.0330[-0.0228, +0.0974]$	$0.429[0.028, 2.094]$
lower 1^+ diagnostic	legacy χ_s (Rohrwild/BBK)	$\theta = 35.3^\circ$	PZ	$+0.337[+0.279, +0.411]$	$14.2[8.93, 21.2]$
lower 1^+ diagnostic	lattice $f_{\gamma,s}^\perp$ (Bacchio et al.)	$\theta = 35.3^\circ$	PZ	$+0.474[+0.420, +0.546]$	$27.8[20.1, 39.4]$

Table 8: Basis/PZ comparison scan for $B_{s1} \rightarrow B_s \gamma$ before the physical-current diagonalization. Entries are medians with 16–84 percentile intervals. These rows are retained as a normalization diagnostic rather than the final physical-current result.

$B_{s1}(5830)$ physical-current decay constants




Figure 12: Monte Carlo width distributions for the observed $B_{s1}(5830) \rightarrow B_s \gamma$ high-combination state after the physical-current diagonalization. The legacy and lattice-normalized photon inputs are shown separately for fixed $\theta = 35.3^\circ$ and for the mixing-angle scan.

The most important lesson from Table 7 is that the observed $B_{s1}(5830)$ amplitude remains cancellation-sensitive after the physical-current normalization is included. The pre-




Figure 13: Central $B_{s1}(5830) \rightarrow B_s \gamma$ stability curves for the high-combination amplitude, width, and unmixed G_A, G_B components.

ferred lattice-normalized row gives a sub-keV median width, while the lower-partner diagnostic remains at the tens-of-keV level. The machine-readable Stage-3 bottom outputs are `outputs/stage3_bs1_physical_decay_constant_scan.csv` and `outputs/stage3_bs1_physical_decay_constant_summary.csv`. The earlier `outputs/stage3_bs1_pz_decay_constant_summary.csv` file is retained as the basis/PZ diagnostic.

21 Combined D_{s1} and B_{s1} Summary

Table 9 collects the current width predictions in one place. The D_{s1} rows include the Stage-3 physical-current residue normalization described above. The B_{s1} rows use Pullin–Zwicky basis-current inputs followed by the mixed-current diagonalization closure summarized in Table 7. In addition to the observed $B_{s1}(5830)$ high-combination state, the table now includes the lower 1^+ bottom-strange diagnostic state used for comparison with the $B_{s1}(5748)$ channel in Bondar–Milstein’s phenomenology. This lower row is not an experimentally established resonance. All entries are medians with 16–84 percentile intervals. The table is intentionally explicit about the photon input, mixing input, normalization status, normalization constants, form factor G , and width Γ , so that no reader has to infer which scenario is being quoted.

Sector	Decay	Photon input	Mixing input	Normalization status	Normalization input	G (GeV $^{-1}$)	Γ (keV)
D_s final	$D_{s1}(2460) \rightarrow D_{s1}\gamma$	legacy χ_s (Rohrwild/BBK)	$\theta = 35.3^\circ$	Stage-3 residue calibrated	$f_1 = 0.344[0.329, 0.363]$, $f_2 = 0.380[0.339, 0.423]$	$-0.424[-0.484, -0.360]$	$37.7^{+11.5}_{-10.5}$
D_s final	$D_{s1}(2536) \rightarrow D_{s1}\gamma$	legacy χ_s (Rohrwild/BBK)	$\theta = 35.3^\circ$	Stage-3 residue calibrated	$f_1 = 0.344[0.329, 0.363]$, $f_2 = 0.380[0.339, 0.423]$	$+0.0118[+0.00143, +0.0233]$	$0.045^{+0.124}_{-0.040}$
D_s final	$D_{s1}(2460) \rightarrow D_{s1}\gamma$	lattice $f_{\gamma,s}^+$ (Bacchio et al.)	$\theta = 35.3^\circ$	Stage-3 residue calibrated	$f_1 = 0.345[0.330, 0.364]$, $f_2 = 0.379[0.338, 0.423]$	$-0.277[-0.312, -0.243]$	$16.1^{+4.4}_{-3.7}$
D_s final	$D_{s1}(2536) \rightarrow D_{s1}\gamma$	lattice $f_{\gamma,s}^+$ (Bacchio et al.)	$\theta = 35.3^\circ$	Stage-3 residue calibrated	$f_1 = 0.345[0.330, 0.364]$, $f_2 = 0.379[0.338, 0.423]$	$+0.0403[+0.0194, +0.0596]$	$0.504^{+0.600}_{-0.377}$
B_s final	$B_{s1}(5830) \rightarrow B_{s1}\gamma$	legacy χ_s (Rohrwild/BBK)	$\theta = 35.3^\circ$	physical-current diagonalized	$f_1 = 0.436[0.380, 0.486]$, $f_2 = 0.218[0.138, 0.292]$	$-0.0793[-0.130, -0.0537]$	$1.33^{+2.27}_{-0.72}$
B_s final	$B_{s1}(5830) \rightarrow B_{s1}\gamma$	lattice $f_{\gamma,s}^+$ (Bacchio et al.)	$\theta = 35.3^\circ$	physical-current diagonalized	$f_1 = 0.443[0.384, 0.492]$, $f_2 = 0.203[0.121, 0.291]$	$-0.0401[-0.0695, -0.0274]$	$0.342^{+0.684}_{-0.182}$
B_s final	$B_{s1}(5830) \rightarrow B_{s1}\gamma$	legacy χ_s (Rohrwild/BBK)	$25^\circ \leq \theta \leq 45^\circ$	physical-current diagonalized	$f_1 = 0.426[0.368, 0.469]$, $f_2 = 0.246[0.174, 0.312]$	$-0.0380[-0.103, +0.0101]$	$0.349^{+1.908}_{-0.323}$
B_s final	$B_{s1}(5830) \rightarrow B_{s1}\gamma$	lattice $f_{\gamma,s}^+$ (Bacchio et al.)	$25^\circ \leq \theta \leq 45^\circ$	physical-current diagonalized	$f_1 = 0.421[0.371, 0.486]$, $f_2 = 0.249[0.170, 0.307]$	$+0.0137[-0.0612, +0.0657]$	$0.493^{+1.233}_{-0.411}$
B_s final	$B_{s1}(5750) \rightarrow B_{s1}\gamma$	legacy χ_s (Rohrwild/BBK)	$\theta = 35.3^\circ$	physical-current diagonalized	$f_1 = 0.428[0.367, 0.479]$, $f_2 = 0.226[0.126, 0.303]$	$+0.284[+0.227, +0.360]$	$9.50^{+3.45}_{-3.37}$
B_s final	$B_{s1}(5750) \rightarrow B_{s1}\gamma$	lattice $f_{\gamma,s}^+$ (Bacchio et al.)	$\theta = 35.3^\circ$	physical-current diagonalized	$f_1 = 0.440[0.381, 0.490]$, $f_2 = 0.216[0.132, 0.285]$	$+0.391[+0.331, +0.457]$	$18.4^{+3.9}_{-3.9}$
B_s final	$B_{s1}(5750) \rightarrow B_{s1}\gamma$	legacy χ_s (Rohrwild/BBK)	$25^\circ \leq \theta \leq 45^\circ$	physical-current diagonalized	$f_1 = 0.418[0.355, 0.475]$, $f_2 = 0.251[0.161, 0.313]$	$+0.295[+0.230, +0.384]$	$10.6^{+7.8}_{-4.3}$
B_s final	$B_{s1}(5750) \rightarrow B_{s1}\gamma$	lattice $f_{\gamma,s}^+$ (Bacchio et al.)	$25^\circ \leq \theta \leq 45^\circ$	physical-current diagonalized	$f_1 = 0.416[0.358, 0.474]$, $f_2 = 0.247[0.160, 0.312]$	$+0.402[+0.339, +0.493]$	$19.6^{+12.3}_{-6.4}$

Table 9: Current combined summary of D_{s1} and B_{s1} radiative form factors and widths. The G column is the E1 coupling in the convention of Eq. (37). The D_s rows include the Stage-3 physical-current residue correction. The B_s rows use Pullin–Zwicky basis-current inputs and the mixed-current diagonalization closure to obtain the physical-current constants f_1 and f_2 . The $B_{s1}(5750)$ rows are lower-partner diagnostics, not observed-state measurements. The older HQ-scaled bottom rows are retained only as diagnostics.

The machine-readable version of Table 9 is `outputs/combined_recommended_results_table.csv`.

21.1 Comparison With Bondar–Milstein Literature Tables

Bondar and Milstein collect several quark-model predictions for the D_{s1} and B_{s1} radiative widths. Their charm table uses Godfrey, Goity–Roberts, Green–Repko–Radford, Radford–Repko–Saelim, Chen–Liu–Zhou–Chen, and Korner–Pirjol–Schilcher for the comparison columns. Their bottom table uses Li–Ni–Zhong, Godfrey–Moats–Swanson, and Lu–Pan–Wang–Wang–Li. We cite these papers explicitly in the manuscript, because otherwise the comparison would look like a citation only to Bondar–Milstein’s summary table.

Table 10 gives a compact version restricted to channels computed in the present LCSR analysis. The comparison set spans early quark-model estimates, relativistic quark-model calculations, more recent potential-model/electric-transition analyses, and the Bondar–Milstein phenomenological extraction. The full machine-readable table also stores the $D_s^*\gamma$ and $B_s^*\gamma$ rows from the literature; our entry is marked as not computed for those channels because they require separate vector-final-state correlators. This scope statement should be kept in the paper so that the comparison table is not read as a claim about uncalculated vector-final-state transitions.

Channel	Early QM	Rel. QM	Recent QM	Bondar–Milstein	This work
$D_{s1}(2536) \rightarrow D_s \gamma$	15 (Godfrey2005)	25.2–31.1 (Goity–Roberts2001)	18.18–18.85 (Chen2020)	12 (Bondar–Milstein2025)	0.504[0.127, 1.104]
$D_{s1}(2460) \rightarrow D_s \gamma$	6.2 (Godfrey2005)	10.3–17.2 (Goity–Roberts2001)	3.53–3.61 (Chen2020)	297 (Bondar–Milstein2025)	16.1[12.4, 20.5]
$B_{s1}(5750) \rightarrow B_s \gamma$	37 (Li–Ni–Zhong2021)	70.6 (Godfrey–Moats–Swanson2016)	97.7 (Lu–Pan–Wang2016)	47 (Bondar–Milstein2025)	18.4[12.5, 28.1]
$B_{s1}(5830) \rightarrow B_s \gamma$	27 (Li–Ni–Zhong2021)	47.8 (Godfrey–Moats–Swanson2016)	56.6 (Lu–Pan–Wang2016)	28 (Bondar–Milstein2025)	0.342[0.160, 1.026]

Table 10: Representative comparison with the radiative-width literature summarized by Bondar–Milstein. Widths are in keV. The entries in the last column use our preferred lattice photon normalization. The complete bookkeeping table is `outputs/bondar_literature_comparison_table.csv`.

22 Comparison With Experimental Information

The sum rule predicts a partial radiative width, $\Gamma_\gamma \equiv \Gamma(H_1 \rightarrow H\gamma)$, while experiments often quote branching fractions, total widths, or spectroscopy measurements in other decay channels. Therefore the comparison is not always one-to-one. In Table 11, the “interpretation” column has the following meaning:

- for $D_{s1}(2460)$, the measured branching fraction is used to infer the total-width scale that would be implied by our partial radiative width;
- for $D_{s1}(2536)$, the measured total width is used to infer the radiative branching fraction implied by our partial width;
- for $B_{s1}(5830)$, no direct radiative measurement is available, so the row is a prediction for future experimental searches rather than a numerical comparison.

Decay	Theory input	Experimental information	Interpretation
$D_{s1}(2460) \rightarrow D_s \gamma$	legacy χ_s : 37.7 [27.2, 49.2] keV	$\mathcal{B}(D_{s1}(2460) \rightarrow D_s \gamma) = (16 \pm 4 \pm 3)\%$ (BaBar)	implies $\Gamma_{\text{tot}} \sim 0.24$ MeV; compatible with a very narrow state
$D_{s1}(2460) \rightarrow D_s \gamma$	lattice $f_{\gamma,s}^+$: 16.1 [12.4, 20.5] keV	$\mathcal{B}(D_{s1}(2460) \rightarrow D_s \gamma) = (16 \pm 4 \pm 3)\%$ (BaBar)	implies $\Gamma_{\text{tot}} \sim 0.10$ MeV; also compatible with a very narrow state
$D_{s1}(2536) \rightarrow D_s \gamma$	legacy χ_s : 0.045 [0.0046, 0.169] keV	$\Gamma_{\text{tot}}(D_{s1}(2536)) = 0.92 \pm 0.05$ MeV (BaBar)	implies $\mathcal{B}_\gamma \sim 0.005\%$; extremely cancellation suppressed
$D_{s1}(2536) \rightarrow D_s \gamma$	lattice $f_{\gamma,s}^+$: 0.504 [0.127, 1.104] keV	$\Gamma_{\text{tot}}(D_{s1}(2536)) = 0.92 \pm 0.05$ MeV (BaBar)	implies $\mathcal{B}_\gamma \sim 0.05\%$; strongly cancellation suppressed
$B_{s1}(5830) \rightarrow B_s \gamma$	legacy χ_s : 1.33 [0.612, 3.60] keV	no direct radiative measurement; observed mainly in $B^{(*)}K$ spectroscopy	prediction for future searches; cancellation-sensitive physical-current amplitude
$B_{s1}(5830) \rightarrow B_s \gamma$	lattice $f_{\gamma,s}^+$: 0.342 [0.160, 1.026] keV	no direct radiative measurement; observed mainly in $B^{(*)}K$ spectroscopy	prediction for future searches; high-combination amplitude remains cancellation sensitive

Table 11: Comparison of the present radiative-width predictions with available experimental information. The BaBar $D_{s1}(2460)$ branching fraction is from the $B \rightarrow D^{(*)}D_{sJ}^{(*)}$ analysis. The $D_{s1}(2536)$ total width is the BaBar width measurement. The $B_{s1}(5830)$ has no direct radiative branching-fraction measurement at present; CMS and LHCb measurements are spectroscopy measurements in hadronic channels.

Numerically, the $D_{s1}(2460)$ rows say that the present partial-width prediction is consistent with a total width well below the MeV scale if the measured radiative branching fraction is used. The $D_{s1}(2536)$ rows say something sharper: with the measured total width near 0.92 MeV, the radiative branching fraction is expected to be far below the percent level. For $B_{s1}(5830)$ the table should be read as a target for future experimental searches, not as a test against an existing radiative measurement.

The machine-readable version of Table 11 is `outputs/experimental_comparison_table.csv`.

23 Paper-Ready Ingredients Before Writing

This section collects the compact version of the calculation choices that should be carried into the manuscript. It is deliberately shorter than the derivation above: the purpose is to state what a reader must know before looking at the final numerical tables.

23.1 Compact Input Table

23.2 Convention Checklist For The Paper

- We use the mostly-minus metric and $\epsilon^{0123} = +1$.

Quantity	Symbol	Value	Comment
charm mass	m_c	1.27 ± 0.02 GeV	$\overline{\text{MS}}$ input for D_s sector; PDG2024
bottom mass	m_b	4.18 ± 0.03 GeV	$\overline{\text{MS}}$ input for B_s sector; PDG2024
strange mass	m_s	0.093 ± 0.011 GeV	retained explicitly, including m_s^2 terms; PDG2024
$D_{s1}(2460)$, $D_{s1}(2536)$, D_s masses	$M_{2460}, M_{2536}, m_{D_s}$	2.4595, 2.53511, 1.96835 GeV	physical charm-strange masses; PDG2024
$B_{s1}(5830)$, B_s masses	M_{5830}, m_{B_s}	5.82870, 5.36692 GeV	observed bottom-strange channel; PDG2024/CMS2018
lower B_{s1} diagnostic mass	$M_{B_{s1}}^{\text{low}}$	5.750 ± 0.026 GeV	expected, not-yet-established lower 1^+ partner; Lang2015
D_s and D_{s1} decay constants	$f_{D_s}, f_{D_{s1}}, f_{D_{s1}}^T$	0.2499, 0.225, 0.256 GeV	hadronic-side inputs; FLAG2024/Colangelo2005/PullinZwicky2021
D_{s1} physical-current residues	f_1, f_2	$0.345 \pm 0.017, \simeq 0.38$ GeV	working Stage-3 anchor; f_2 from overlap calibration; checked against Wang2015 pure axial benchmark
B_s and working B_{s1} decay constants	$f_{B_s}, f_{B_{s1}}, f_{B_{s1}}^T$	0.2303, 0.146, 0.164 GeV	diagnostic HQ-scaled inputs, retained only as a cross-check
PZ B_{s1} basis-current inputs	$f_{B_{s1}}^{PZ}, f_{B_{s1}}^{PZ,T}$	$0.305 \pm 0.027, 0.285 \pm 0.027$ GeV	dedicated bottom-sector basis inputs; PullinZwicky2021
B_{s1} physical-current constants	$f_1^{(B_{s1})}, f_2^{(B_{s1})}$	$0.443[0.384, 0.492], 0.203[0.121, 0.291]$ GeV	fixed-angle lattice scenario from mixed-current diagonalization closure
condensates	$\langle \bar{q}q \rangle, \kappa$	$-(0.240 \text{ GeV})^3, 0.8 \pm 0.1$	$\langle \bar{s}s \rangle = \kappa \langle \bar{q}q \rangle$; FLAG2024/standard SVZ input
legacy photon input	$\chi_s(1 \text{ GeV})$	$+3.15 \pm 0.30 \text{ GeV}^{-2}$	Rohrwild/BBK-style comparison scenario; BBK2002/Rohrwild2007
lattice photon input	$f_{\gamma,s}^+(1 \text{ GeV})$	$-55.1 \pm 1.9 \text{ MeV}$	preferred leading-twist normalization scenario; Bacchio2024
mixing angle	θ	35.3° central, 25° – 45° scan	3P_1 – 1P_1 physical-state mixing; Colangelo2005 convention
D_s sum-rule window	$M^2, \sqrt{s_0}$	$3\text{--}6 \text{ GeV}^2, 2.50\text{--}2.60 \text{ GeV}$	charm-sector stability/uncertainty scan
B_s sum-rule window	$M^2, \sqrt{s_0}$	$8\text{--}14 \text{ GeV}^2, 6.05\text{--}6.25 \text{ GeV}$	observed $B_{s1}(5830)$ scan

Table 12: Compact input table for the manuscript. The complete machine-readable version is `outputs/paper_input_summary_table.csv`; the longer bookkeeping table is `inputs_table.csv`.

- The initial axial-meson momentum is $P = p + q$; p is the outgoing pseudoscalar momentum and q is the photon momentum.
- The tensor current is $i\bar{s}\sigma_{\mu\alpha}P^\alpha\gamma_5 Q/(m_Q + m_s)$, so the momentum entering the tensor current is the initial-state momentum P .
- The physical amplitudes are $G_{\text{low}} = \sin\theta G_A + \cos\theta G_B$ and $G_{\text{high}} = \cos\theta G_A - \sin\theta G_B$ in the basis-current normalization convention.
- For the physical currents J_1 and J_2 , the final couplings should be divided by the corresponding two-point residues f_1 and f_2 as in Eqs. (62)–(63).
- The width is computed from $\Gamma = \alpha_{\text{em}} G^2 |\vec{q}_\gamma|^3/3$.
- We store $\chi_s(1 \text{ GeV})$ as a positive Rohrwild/BBK magnitude, $\chi_s = +3.15 \text{ GeV}^{-2}$; this must not be mixed with papers that print the opposite sign convention.

23.3 Propagator And Condensate Statement

The base propagators in the hard calculation are not expanded in ordinary vacuum condensates. Instead, we use the quark propagators in an external electromagnetic field for hard photon emission and the background-gluon propagator term for the three-particle photon DAs. The condensates that appear in the final formulas are therefore part of the photon matrix elements and photon DA normalizations, such as $\chi_s \langle \bar{s}s \rangle$, not separate SVZ condensate corrections to the free propagator. This distinction should be stated clearly in the paper because it is one of the main convention choices of the calculation.

23.4 Quoting Strategy

For the final paper tables, the lattice-normalized $f_{\gamma,s}^\perp$ scenario should be presented as the preferred modern input choice, while the legacy $\chi_s \langle \bar{s}s \rangle$ scenario should be kept for comparison with the older photon-DA literature. Uncertainties should be quoted as medians with 16–84 percentile intervals. This is essential for $D_{s1}(2536)$ and $B_{s1}(5830)$ because their amplitudes can pass close to zero, making Gaussian errors misleading.

For the bottom-strange sector, the observed $B_{s1}(5830)$ prediction is a future-search target. The lower B_{s1} row is a prediction for an expected but not yet experimentally established 1^+ state; it should not be described as an observed resonance. Its role in the manuscript is to show that the lower combination is constructive and can have a much larger radiative width than the observed high-combination state.

23.5 Reference Map

The manuscript reference list should connect each input family to its source: PDG for masses and quark charges, FLAG Review 2024 for f_{B_s} and related lattice decay-constant context, Rohrwild/BBK for the legacy photon DAs, Bacchio et al. for the lattice $f_{\gamma,s}^\perp$ normalization, Colangelo, De Fazio and Ozpineci for the original charm-strange LCSR benchmark, Pullin and Zwicky for heavy-light decay-constant cross-checks, Lang–Mohler–Prelovsek–Woloshyn for the lower B_{s1} diagnostic mass, BaBar for the $D_{s1}(2460)$ branching fraction and $D_{s1}(2536)$ total width, and CMS/LHCb for the observed $B_{s1}(5830)$ spectroscopy. The Bondar and Bondar–Milstein papers should be cited in the discussion of the $D_{s1}(2536)$ nonobservation, the constructive/destructive E1-amplitude mechanism, and the comparison with the lower and higher B_{s1} radiative channels.

Citation key	Input family	Used for
PDG2024	masses, quark masses, charges, α_{em}	kinematics, quark-mass inputs, width normalization
FLAG2024	lattice decay constants and condensate context	f_{D_s} , f_{B_s} , and modern lattice normalization context
Colangelo2005	charm-strange radiative LCSR benchmark	comparison calculation and charm axial-current normalization
BBK2002	photon distribution amplitudes	twist-2, twist-3, and twist-4 photon DA shapes
Rohrwild2007	magnetic susceptibility update	legacy χ_s comparison scenario
Bacchio2024	lattice photon-DA normalization	preferred $f_{\gamma,s}^\perp(1 \text{ GeV})$ scenario
PullinZwicky2021	heavy-light vector/tensor decay constants	D_s tensor scale cross-check and preferred B_{s1} PZ scenario
Wang2015	heavy-light decay constants from QCDSR	pure axial-current $f_A(D_{s1})$ benchmark for the Stage-3 residue check
Lang2015	positive-parity B_s spectroscopy	lower B_{s1} diagnostic mass
BaBar2006	$D_{s1}(2460)$ branching data	experimental interpretation of $\Gamma(D_{s1}(2460) \rightarrow D_s \gamma)$
BaBar2011	$D_{s1}(2536)$ total width	branching-fraction estimate for the high charm-strange state
CMS2018	$B_{s1}(5830)$ spectroscopy	observed bottom-strange mass and phenomenological comparison
Bondar2025Why	$D_{s1}(2536)$ radiative nonobservation	experimental upper-limit motivation for the suppressed high-state mode
BondarMilstein2025	D_{s1} and B_{s1} radiative phenomenology	constructive/destructive E1-amplitude mechanism and comparison tables
Godfrey2005, GoityRoberts2001, Chen2020	D_{s1} radiative-width literature	comparison columns for $D_{s1} \rightarrow D_s^{(*)} \gamma$
LiNiZhong2021, GodfreyMoatsSwanson2016, LuPanWang2016	B_{s1} radiative-width literature	comparison columns for $B_{s1} \rightarrow B_s^{(*)} \gamma$

Table 13: Citation map for the numerical inputs and comparison data. The machine-readable version is `outputs/input_citation_map.csv`.

24 Pre-Paper Calculation Closure Checklist

The core D_{s1} and B_{s1} calculation set is complete at the precision level used in these working notes. The completed items are:

- Stage 1 hard-photon and two-particle photon-DA terms, with both heavy- and strange-line hard emission included;
- Stage 2 three-particle photon-DA corrections from the background-gluon insertion in the quark propagator;
- explicit retention of the m_s and m_s^2 terms used by the present Mathematica/Python kernels;
- D_{s1} Stage 3 physical-current normalization using the calibrated residues $f_1 \simeq 0.345 \text{ GeV}$ and $f_2 \simeq 0.38 \text{ GeV}$, while using Wang2015 only as a pure-axial benchmark rather than as a direct authority for the mixed residues;
- B_{s1} Stage 3 normalization with Pullin–Zwicky basis-current inputs and the mixed-current diagonalization closure, plus the lower-state diagnostic channel;
- one-at-a-time uncertainty diagnostics, Monte Carlo uncertainty scans, percentile intervals, stability plots in (M^2, s_0) , and comparison tables;
- experimental interpretation tables for the available D_{s1} and B_{s1} information;
- citation-tagged input tables in `inputs_table.csv`, `outputs/paper_input_summary_table.csv`, and `outputs/input_citation_map.csv`.

The only non-blocking physics refinement left before submission is an independent local two-point validation of the complete current matrix:

- derive the local two-point OPE for Π_{AA} , Π_{BB} and Π_{AB} , including the ordinary SVZ condensates appropriate to a two-point decay-constant sum rule;
- compare the extracted f_1 and f_2 with the Stage-3 calibrated values, $f_1 \simeq 0.345$ GeV and $f_2 \simeq 0.38$ GeV;
- rerun the D_{s1} headline table only if the independent two-point residues differ appreciably from the calibrated values;
- optionally repeat the same fully mixed two-point check in the bottom-strange sector once a dedicated local-QCDSR normalization is desired.

Private future-work note. The B_{c1} radiative system should not be included in the present paper. It is a heavy-heavy problem rather than a heavy-light strange-meson problem, so the present photon-DA LCSR machinery does not carry over without a new formalism. We keep this only as an internal next-study idea in `notes/internal_future_work.md`.

25 Redo Analysis and Publication-Quality Plots

From this point onward these notes are the active calculation document. The folder `draft_prd/` is kept only as a reference copy of the fuller PRD-style derivation, and the folder `manuscript/` is no longer the editing target for the paper. The paper text will be edited elsewhere; here we will keep the reproducible analysis, numerical checks, and plot-generation record.

The next analysis pass should be done conservatively. Existing output files are not deleted; revised scans should be written with clear names so that old and new results can be compared. The main goals are:

- recheck the final D_{s1} and B_{s1} numerical inputs, especially the photon normalization, physical-current decay constants, Borel windows, continuum thresholds, and mixing-angle treatment;
- rerun the preferred Monte Carlo scans with an explicit random seed and store the seed, accepted sample count, rejected sample count, and percentile definitions in the output summary;
- separate the fixed-angle $\theta = 35.3^\circ$ results from the $25^\circ \leq \theta \leq 45^\circ$ scan in both tables and plots;
- keep the legacy Rohrwild/BBK susceptibility rows as comparison rows, but make the lattice photon-normalized rows the headline results;
- verify the stability windows by producing publication-quality (M^2, s_0) stability plots for the final quoted channels;
- produce compact uncertainty-budget plots showing the dominant input sources, with the photon normalization, condensate combination, Borel window, continuum threshold, and mixing angle varied in a transparent way;
- produce final comparison tables and figures against the representative D_{s1} and B_{s1} literature values, while clearly marking channels not computed in this LCSR setup;
- record every script command used to regenerate the headline figures and tables, so that the paper plots can be reproduced without relying on memory.

Publication figures should be generated as both PDF and PNG. The PDF is used for the paper, while the PNG is useful for quick inspection. The preferred plot style should be consistent across all figures: readable axis labels, explicit units, no crowded legends, and enough whitespace for journal-column scaling. The working output directory remains `outputs/`; if a figure becomes a final paper figure, its source script and source data should be listed in `outputs/manuscript_reproducibility_map.csv`.

Before any revised result replaces a number in the paper, the following checks should be recorded in this section or in a linked table:

- central-value reproduction from the script output;
- percentile interval definition;
- comparison with the previous stored result;
- stability under the chosen Borel and threshold ranges;
- sign convention and current-normalization convention used for the quoted form factor.

25.1 Redo Borel and Continuum Windows

In this redo pass we use the notation s_0 for the squared continuum threshold, in GeV^2 . When it is helpful to compare with the older project scripts, we also give the corresponding value of $\sqrt{s_0}$ in prose, but the tables and new plots use s_0 . This avoids the dimensional ambiguity that would arise from writing, for example, $s_0 = 2.85 \text{ GeV}$.

The first broad stability scan is generated by

```
python3 scripts/redo_stability_windows.py
```

and the output files are

- `outputs/redo_ds1_stability_windows.csv`, `outputs/redo_ds1_stability_windows.pdf`;
- `outputs/redo_bs1_stability_windows.csv`, `outputs/redo_bs1_low_stability_windows.pdf`;
- `outputs/redo_bs1_5830_stability_windows.pdf`.

The central photon input in this diagnostic is the lattice $f_{\gamma,s}^\perp$ normalization. The legacy Rohrwild/BBK susceptibility normalization should be kept as a comparison scenario, but it is not used for choosing the preferred central stability plots.

Table 14: Broad Borel and continuum windows used in the redo stability scan. The s_0 entries are squared thresholds. The $\sqrt{s_0}$ column is only shown to make the connection with the older scripts transparent.

Channel	M^2 broad scan (GeV^2)	s_0 broad scan (GeV^2)	corresponding $\sqrt{s_0}$ (GeV)
$D_{s1}(2460), D_{s1}(2536)$	2.0 – 7.0	8.12 – 9.61	2.85 – 3.10
$B_{s1}(5750)$	6.0 – 18.0	36.60 – 40.32	6.05 – 6.35
$B_{s1}(5830)$	6.0 – 18.0	37.82 – 40.96	6.15 – 6.40

The first conclusion from these plots is methodological rather than numerical. For the bottom-strange observed channel, the old working interval $8 \leq M^2 \leq 14 \text{ GeV}^2$ is reasonable as a conservative window, and the slightly shifted interval $10 \leq M^2 \leq 16 \text{ GeV}^2$ is also defensible from the shape of Fig. 15. For the charm-strange channel, however, the larger above-pole threshold range $s_0 = 8.12 - 9.61 \text{ GeV}^2$ produces a sizeable drift in both M^2 and s_0 . Therefore the final D_{s1} uncertainty scan should not yet be rerun with this full broad threshold interval. The next check should quantify pole dominance and OPE convergence for the old near-threshold window and the new above-pole window side by side, then choose the quoted charm-sector window from the overlap of those criteria.




Figure 14: Redo $D_{s1} \rightarrow D_s \gamma$ Borel stability curves using s_0 as the squared threshold. The broad $s_0 = 8.12 - 9.61 \text{ GeV}^2$ scan is useful diagnostically, but the D_{s1} curves are not flat over the full range. This means the final charm-sector window should be fixed only after applying the pole-dominance and OPE-convergence checks, not simply by taking the whole broad scan as the uncertainty range.

25.2 Pole Proxy, OPE Hierarchy, and Final Window Choice

The follow-up diagnostic is generated by

```
python3 scripts/window_diagnostics.py
```

with outputs `outputs/window_diagnostics_grid.csv`, `outputs/window_diagnostics_summary.csv`, `outputs/window_diagnostics_summary.txt`, and `outputs/window_diagnostics_summary.pdf`.

The pole-dominance entry in this table is an absolute perturbative spectral-weight proxy. This wording is important: the hard spectral terms have explicit dispersion integrals, while the photon-DA contact terms do not all have the same continuum representation. The proxy is therefore used to compare windows, not as an independent physical observable.

The preferred working windows for the final Monte Carlo pass are therefore

$$D_{s1} : \quad 3.0 \leq M^2 \leq 4.5 \text{ GeV}^2, \quad 7.84 \leq s_0 \leq 8.70 \text{ GeV}^2,$$

$$B_{s1}(5750) : \quad 10 \leq M^2 \leq 14 \text{ GeV}^2, \quad 36.60 \leq s_0 \leq 40.32 \text{ GeV}^2,$$

and

$$B_{s1}(5830) : \quad 10 \leq M^2 \leq 14 \text{ GeV}^2, \quad 37.82 \leq s_0 \leq 40.96 \text{ GeV}^2.$$

For D_{s1} , the old near-threshold window had acceptable apparent stability but a poor pole proxy. The broad above-pole window improved the threshold placement but was too unstable. The preferred window is the compromise: it keeps the axial and tensor spectral proxies around or




Figure 15: Redo $B_{s1}(5830) \rightarrow B_s \gamma$ stability curves. The broad scan shows a noticeably flatter region for the quoted high-combination amplitude around $M^2 \simeq 10 - 16 \text{ GeV}^2$, with a moderate residual threshold dependence.

above 0.5, keeps the three-particle terms small, and does not inflate the $D_{s1}(2460)$ stability span relative to the old window. The $D_{s1}(2536)$ span remains large because this channel is controlled by a near cancellation in the physical amplitude; this is a physical sensitivity, not by itself a failure of the Borel window.

For B_{s1} , the narrowed interval $10 \leq M^2 \leq 14 \text{ GeV}^2$ is preferred over $8 \leq M^2 \leq 14 \text{ GeV}^2$ if the main criterion is the stability of the quoted width. However, lowering the Borel interval improves the pole proxy, so we performed a focused bottom-sector scan with fixed continuum thresholds and candidate Borel windows. The script

```
python3 scripts/bs1_window_candidate_scan.py
```

writes `outputs/bs1_borel_window_candidate_summary.csv` and `outputs/bs1_borel_window_candidate_s`. The compact result is

State	$M^2 \text{ (GeV}^2\text{)}$	pole A	pole B	width span	max 3p
$B_{s1}(5750)$	8 – 12	0.456	0.705	0.641	0.0010
$B_{s1}(5750)$	9 – 13	0.405	0.662	0.544	0.0010
$B_{s1}(5750)$	10 – 14	0.360	0.618	0.478	0.0009
$B_{s1}(5830)$	8 – 12	0.475	0.736	0.415	0.0010
$B_{s1}(5830)$	9 – 13	0.423	0.693	0.292	0.0009
$B_{s1}(5830)$	10 – 14	0.378	0.648	0.238	0.0009

Thus lowering M^2 does move the axial pole proxy closer to the usual 0.5 guideline, especially for $8 \leq M^2 \leq 12 \text{ GeV}^2$, but the price is a larger width span. The three-particle fractions remain




Figure 16: Redo $B_{s1}(5750) \rightarrow B_s \gamma$ diagnostic stability curves. The lower-state prediction has a stronger monotonic M^2 dependence than the observed $B_{s1}(5830)$, so its quoted range should remain a model diagnostic until the corresponding state assignment and continuum window are better fixed.

negligible in all these bottom-sector windows. Therefore the choice is a controlled tradeoff: 10 – 14 is the most stable window, while 9 – 13 can be used as a reasonable systematic cross-check because it improves the pole proxy without the full instability of 8 – 12.

25.3 Publication Stability Figures

For the manuscript, the cleanest stability demonstration is not a side-by-side plot. We therefore generate the two standard checks as separate figures: one figure varies M^2 at selected fixed s_0 , and a second figure varies s_0 at selected fixed M^2 . This gives more control over the vertical scale and makes the stability of each physical channel easier to judge. The script

```
python3 scripts/publication_stability_selected_windows.py
```

also writes combined diagnostic plots and the numerical grid:

- outputs/publication_selected_window_stability.pdf;
- outputs/publication_selected_window_stability_normalized.pdf;
- outputs/publication_selected_window_stability.csv.

The combined normalized version is useful diagnostically, but it is not the cleanest publication figure because solid/dashed curve conventions and multiple channels must be decoded simultaneously. For the paper draft, use the channel-separated single-panel figures instead:

- outputs/publication_stability_ds1_2460_M2.pdf, outputs/publication_stability_ds1_2460_s0

Table 15: Window diagnostics for the final redo choice. The pole entries are minimum absolute perturbative spectral-weight fractions over the window. The stability spans are relative spans over the scan grid.

Window	M^2 (GeV ²)	s_0 (GeV ²)	pole A min	pole B min	main span	max 3p
D_{s1} old near-threshold	3.0 – 6.0	6.25 – 6.76	0.280	0.378	0.694	0.016
D_{s1} broad above-pole	2.0 – 7.0	8.12 – 9.61	0.364	0.490	1.485	0.029
D_{s1} preferred	3.0 – 4.5	7.84 – 8.70	0.523	0.620	0.628	0.008
$B_{s1}(5750)$ old diagnostic	8.0 – 14.0	36.60 – 40.32	0.360	0.618	0.765	0.001
$B_{s1}(5750)$ preferred diagnostic	10.0 – 14.0	36.60 – 40.32	0.360	0.618	0.478	0.001
$B_{s1}(5830)$ old diagnostic	8.0 – 14.0	37.82 – 40.96	0.378	0.648	0.423	0.001
$B_{s1}(5830)$ preferred	10.0 – 14.0	37.82 – 40.96	0.378	0.648	0.238	0.001

- `outputs/publication_stability_ds1_2536_M2.pdf`, `outputs/publication_stability_ds1_2536_sc`
- `outputs/publication_stability_bs1_low_5750_M2.pdf`, `outputs/publication_stability_bs1_low`
- `outputs/publication_stability_bs1_5830_M2.pdf`, `outputs/publication_stability_bs1_5830_sc`

For the manuscript, one representative pair is likely enough in the main numerical section. The $D_{s1}(2460) \rightarrow D_s \gamma$ pair is a good default because it is the least cancellation-dominated charm-strange channel. The remaining channel-separated figures should be kept in the project files and can be mentioned as reproducibility checks or moved to supplemental material if the journal format allows it.

25.4 Final-Window Monte Carlo Uncertainty

After fixing the working windows, the final uncertainty should be assigned from a Monte Carlo scan rather than from the visual spread of the stability curves. The stability curves justify the admissible M^2 and s_0 domain; the Monte Carlo scan then propagates the input parameters and the residual window dependence inside that domain. The dedicated final pass is generated by

```
python3 scripts/final_window_monte_carlo.py
```

with fixed random seed 20260705 and 500 attempted samples per state/scenario/angle ensemble/window. The script writes `outputs/final_window_mc_scan.csv`, `outputs/final_window_mc_summary.csv`, `outputs/final_window_mc_summary.txt`, `outputs/final_window_mc_width_histograms.pdf`, and `outputs/final_window_bs_window_crosscheck.pdf`.

The preferred photon input remains the lattice $f_{\gamma,s}^\perp$ normalization. The recommended presentation is the median with the 16–84 percentile interval, especially for $D_{s1}(2536)$ and $B_{s1}(5830)$, where cancellations make the width distribution non-Gaussian near zero. For the angle-scan ensemble and the central windows, the preferred rows are

Channel	g (GeV ⁻¹)	Γ (keV)
$D_{s1}(2460) \rightarrow D_s \gamma$	-0.414[-0.453, -0.376]	36.1[29.7, 43.1]
$D_{s1}(2536) \rightarrow D_s \gamma$	+0.030[-0.014, +0.075]	0.422[0.026, 1.760]
$B_{s1}(5750) \rightarrow B_s \gamma$	+0.394[+0.350, +0.462]	19.3[14.2, 27.2]
$B_{s1}(5830) \rightarrow B_s \gamma$	+0.010[-0.079, +0.060]	0.498[0.056, 1.680]

These values supersede the earlier uncertainty tables if the final windows in Table 15 are adopted. In particular, the above-pole D_{s1} threshold window increases the preferred $D_{s1}(2460)$ width relative to the older near-threshold Monte Carlo pass. This shift is a useful warning: the




Figure 17: Compact comparison of candidate windows. The dashed line in the upper panel marks a pole-fraction proxy of 0.5; values below this level are treated as warnings rather than hard exclusions because only the perturbative spectral part is used in this proxy.

continuum-window choice is a real systematic effect and should be discussed transparently in the manuscript.

For the bottom sector, the final scan also compares the central $10 \leq M^2 \leq 14 \text{ GeV}^2$ interval with the $9 \leq M^2 \leq 13 \text{ GeV}^2$ cross-check window, while keeping the same s_0 intervals. The $B_{s1}(5750)$ median changes from 19.3 to 21.1 keV, and the $B_{s1}(5830)$ median changes from 0.498 to 0.524 keV in the lattice angle-scan ensemble. This is small compared with the quoted percentile widths, so the central bottom-sector window is stable against this reasonable Borel-window variation.




Figure 18: M^2 -dependence of the physical form factor for $D_{s1}(2460) \rightarrow D_s \gamma$, shown at selected fixed s_0 values. The vertical range is deliberately kept wider than the curve spread so that the visual impression reflects the actual stability of the chosen window.




Figure 19: s_0 -dependence of the physical form factor for $D_{s1}(2460) \rightarrow D_s \gamma$, shown at selected fixed M^2 values. This is the complementary threshold-stability check for the same selected window.




Figure 20: Final-window Monte Carlo width distributions for the preferred lattice photon input and $25^\circ \leq \theta \leq 45^\circ$. Solid vertical lines mark the median and dashed lines mark the 16–84 percentile interval.




Figure 21: Bottom-sector final-window Monte Carlo cross-check. The $9 \leq M^2 \leq 13 \text{ GeV}^2$ interval improves the pole proxy but gives results compatible with the central $10 \leq M^2 \leq 14 \text{ GeV}^2$ interval within the final percentile uncertainty.